Philipp Hilbert. *Product Manager who turns discovery research into shipped search experiences.*

Positioning, ten years of B2B SaaS product ownership — including document management, data pipelines, and AI-assisted workflows — applied to Preservica's hardest unsolved problem: helping archivists find what they need.

BERLIN, GERMANY · HILBERT@TRUE-NORTH.BERLIN · TRUE-NORTH.BERLIN

LinkedIn · GitHub · www.true-north.berlin · YouTube (EN sample)

§ WHY PRESERVICA

The framing in the job description is precise: archivists have historically had to know exactly what they are looking for, and AI is changing that at a structural level. That is not a vague AI trend; it is a concrete shift in what the product can promise its users. The work of translating semantic search and natural language querying into a backlog grounded in real archivist workflows — rather than in capability demos — is exactly the kind of problem I have been solving across regulated, data-heavy environments for the past decade. Preservica's customer base includes institutions with decades of inconsistent cataloguing practice, which means the discovery challenge is genuinely hard and the research cannot be shortcut. *That is precisely where careful discovery work and risk-first prioritisation compound.*

§ HONEST FRAMING

I have not worked directly in digital preservation or archival science as an industry vertical. What I do bring from document management implementations at projuncta and CLINET, data provenance and auditability work at Rohde & Schwarz, and ETL pipeline ownership at the Federal Employment Agency is a strong mechanic for thinking about how heterogeneous, incomplete, and historically inconsistent data gets surfaced reliably — which is the underlying engineering and product challenge here.

§ AI-NATIVE STACK

Daily tools, not buzzwords.

- **Claude**, specification drafting, edge-case generation, research synthesis
- **GPT**, user story refinement and stakeholder communication drafts
- **GitHub Copilot**, in-repo review and acceptance criteria validation
- **Cursor / VSCode**, reading and navigating production codebases
- **Lovable**, rapid front-end prototyping before engineering commitment
- **Framer**, design-heavy interactive prototypes for user testing
- **Gamma.app**, stakeholder-facing roadmap and research readouts
- **n8n**, agentic workflow design and multi-step automation

§ TECHNICAL ENVIRONMENT

Shipped with, not just listed.

Python, Django, JavaScript, PostgreSQL, MongoDB, Kotlin (Spring Boot), Redpanda, Dagster, dbt, DuckDB, Kubernetes, ArgoCD, Helm, Argo Workflows, OpenShift, Jenkins, Bitbucket, REST APIs, AngularJS, Ionic, MicroStrategy, GitOps pipelines.

§ SIDE PROJECT

BACKTEST ENGINE & DEV AGENT FRAMEWORK GITHUB.COM/HILBERTP →

Personal infrastructure project: a crypto backtesting engine and autonomous dev-agent framework built from scratch.

- Implements a file-queue-based agent loop for autonomous code generation and self-correction across iterative builds.
- Architecture uses Python with modular strategy interfaces, separating signal generation from execution simulation cleanly.
- Actively maintained across multiple repositories; serves as a live testbed for agentic workflows and LLM-assisted development patterns.
- Demonstrates the same self-directed, technically grounded product execution style that underpins backlog ownership at Preservica — where AI capability must be understood at implementation depth, not just described in requirements.

§ HOW I WOULD WORK AT PRESERVICA

First weeks, concrete.

WEEK 1

Listen, map, find the gaps.

Shadow customer-facing teams, read existing research and support tickets, and map how archivists currently navigate the Preservica interface to find content — paying specific attention to where users abandon a search or escalate to specialist colleagues.

WEEK 2

Frame, slice, validate.

Synthesise findings into three to five ranked problem statements, bring them back to the Tech Lead and ADM for a shared feasibility read, and identify the smallest prototype that would let real users confirm or refute the most uncertain assumption.

WEEK 3+

Ship, document, compound.

Begin running structured user research sessions alongside early prototypes, feed findings into the backlog with clear acceptance criteria, and establish a shared research log so the team's discovery work accumulates rather than resets sprint to sprint.

I COVER LETTER

Dear Preservica hiring team,

The problem you describe is one I find genuinely interesting: archivists working with large, heterogeneous collections currently have to know the right terminology and the correct metadata reference to find anything. Semantic search and natural language querying can change that, but only if the product work is grounded in how archivists actually search — not in what the technology can do in a demo. That distinction between a capability and a solved user problem is the exact tension where I have spent most of my product career.

At Rohde & Schwarz I owned a central AI and data management platform built for multimodal datasets under the FCAS defence programme. The core challenge was not the ML capability itself but building the data provenance, audit trails, and filtering pipelines that made model outputs trustworthy to the scientists and engineers consuming them. That experience taught me how AI-powered features introduce risk in ways that pure retrieval systems do not — and that managing that risk requires careful requirements work before engineering effort is committed, not after. At the Federal Employment Agency I owned a public analytics platform where the underlying data was regional, historically inconsistent, and politically sensitive. Designing search and filter behaviour for users who did not always know the right administrative terminology – and who needed to trust the results – maps directly onto the archivist workflow problem Preservica is solving.

Earlier in my career I implemented document management systems both at my own consultancy and for CLINET Platforms, and I worked as Proxy Product Owner on the migration of a complex internal catalogue at AOK health insurance. These were not glamorous projects, but they gave me a concrete understanding of how document collections degrade over time, how metadata becomes inconsistent across organisational handovers, and how users develop workarounds when search fails them. That context makes me a more useful partner to the engineering team when trade-offs between indexing strategy, relevance ranking, and UI design need to be made with real user behaviour in mind.

I work AI-native by default: I use Claude and GPT for specification drafting and edge-case generation, Lovable and Framer for rapid prototyping before engineering commitment, and I have built my own agentic workflow tooling in n8n. I understand AI capabilities and limitations at a level that lets me make sensible product decisions about AI-powered features – including the contexts where they introduce risk or need careful design guardrails. Shipping AI features that archivists and records managers can actually trust is a harder problem than shipping AI features that impress in a sales demo, and I take that distinction seriously.

I am based in Berlin, authorised to work in the UK remotely as an EU citizen, and available within four weeks of an offer.

Best regards, *Philipp Hilbert*

II CURRICULUM VITAE

---

§ BEARING

Product Manager with nearly ten years of B2B SaaS experience across AI data platforms, public analytics, regulated insurance technology, and document management. Comfortable in environments where requirements are incomplete and discovery cannot be shortcut: the work at Rohde & Schwarz required building trust in AI-generated outputs through provenance and auditability; the work at the Federal Employment Agency required designing search and filter behaviour for heterogeneous, historically inconsistent data. Both map directly onto the challenge of making Preservica's preserved collections discoverable to users who do not always know the right terminology. Strong delivery ownership, risk-first discovery framing, and AI-native working practice.

---

§ CORE STRENGTHS

- Risk-first discovery: value, usability, feasibility, and viability framed before engineering starts
- Translating messy, heterogeneous data problems into dev-ready requirements
- AI-native product execution: prototyping, spec drafting, and edge-case generation
- Data provenance, audit trails, and trustworthy AI feature design
- Backlog ownership and sprint discipline in agile delivery teams
- Cross-functional alignment across engineering, data science, QA, and compliance stakeholders
- Document and content management system implementation experience
- ETL pipeline and data platform product ownership
- Founder-level pragmatism: low ego, high agency, comfortable with ambiguity

---

§ PROFESSIONAL EXPERIENCE

ROHDE & SCHWARZ 2024, 2025

Product Manager, AI Data Transformation

- Owned data provenance, audit trails, and reproducible transformation pipelines for multimodal AI training datasets — the same trust and traceability challenges that underpin trustworthy search results in a preservation context
- Designed SafeAI and ExplainableAI evaluation concepts to make model outputs interpretable to non-ML stakeholders
- Delivered autonomous feature sets end-to-end: user stories, acceptance criteria, implementation support, testing, and GitOps rollout
- Reduced deployment times from several days to 20–40 minutes by introducing GitOps with ArgoCD
- Led a cross-functional team of architect, data scientists, backend and frontend engineers, and QA

Tech: Kotlin (Spring Boot), Python, PostgreSQL, MongoDB, Redpanda, Kubernetes, Argo Workflows, ArgoCD, OpenShift

---

FEDERAL EMPLOYMENT AGENCY (BUNDESAGENTUR FÜR ARBEIT) 2023, 2024

Product Manager, Public Data Analytics

- Designed search, filter, and comparison features for regional labour market data that users without specialist terminology could navigate confidently
- Introduced Dagster for automated ETL and indicator generation across historically inconsistent regional datasets
- Served as Deputy IT Security Officer with focus on compliance, governance, and data lineage
- Migrated legacy systems to Kubernetes and standardised CI/CD with full test integration

Tech: Python, Django, JavaScript, PostgreSQL, dbt, Dagster, DuckDB, MicroStrategy, Kubernetes, Jenkins

---

EMIL GROUP 2022

Product Manager, InsurTech SaaS Platform

- Delivered document management, workflow automation, and policy issuance modules on a B2B platform serving primary insurers, reinsurers, MGAs, and brokers
- Designed and shipped a Claims Center module from concept to production rollout
- Re-established sprint discipline by splitting a bloated 17-person team into focused Scrum units with clear ownership
- Regained client trust in a prototype-stage environment through structured communication and weekly product workshops

Tech: Low-code platform, Java, HTML/CSS, Kubernetes, REST APIs

---

CLINET PLATFORMS 2021, 2022

Product Owner, Mobile Healthcare App

- Built patient document storage and retrieval flows as part of a mobile-first hospital app, including integration with legacy hospital information system APIs
- Delivered core flows for digital anamnesis, document management, and clinical data access under tight deadlines
- Communicated technical complexity of legacy API constraints transparently to non-technical stakeholders

Tech: AngularJS, Ionic, Python, Kubernetes, CGM API integration

---

§ FOUNDER & EARLY-STAGE EXPERIENCE

QCRYPT AG 2016, 2018

Product Owner & Deputy IT Security Officer

- Owned product architecture for a quantum-secure encryption product: backup, restore/checkpoint, and data encryption flows
- Designed for non-technical users: full setup under two minutes with no prior technical knowledge required
- Served as deputy IT security officer; responsible for attack vector mitigation and security governance

KVITT PAYMENT SOLUTIONS 2013, 2018

Founder & CFO

- Co-founded a group-based P2P payments product; owned PSP integration, full UX/UI design, and mobile app prototype
- Later acquired and operated by Sparkasse

AOK HEALTH INSURANCE – ITSCARE 2019, 2020

Proxy Product Owner

- Migrated internal webshop to Shopware, including complex catalogue and API integration and customisable approval policies

---

§ LANGUAGES

German, native. English, C2.

LinkedIn · GitHub · www.true-north.berlin · YouTube (EN sample)
