

Philipp Hilbert. *Product Manager who turns discovery research into shipped search experiences.*

Positioning, ten years of B2B SaaS product ownership – including document management, data pipelines, and AI-assisted workflows – applied to Preservica's hardest unsolved problem: helping archivists find what they need.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear Preservica hiring team,

The problem you describe is one I find genuinely interesting: archivists working with large, heterogeneous collections currently have to know the right terminology and the correct metadata reference to find anything. Semantic search and natural language querying can change that, but only if the product work is grounded in how archivists actually search – not in what the technology can do in a demo. That distinction between a capability and a solved user problem is the exact tension where I have spent most of my product career.

At Rohde & Schwarz I owned a central AI and data management platform built for multimodal datasets under the FCAS defence programme. The core challenge was not the ML capability itself but building the data provenance, audit trails, and filtering pipelines that made model outputs trustworthy to the scientists and engineers consuming them. That experience taught me how AI-powered features introduce risk in ways that pure retrieval systems do not – and that managing that risk requires careful requirements work before engineering effort is committed, not after. At the Federal Employment Agency I owned a public analytics platform where the underlying data was regional, historically inconsistent, and politically sensitive. Designing search and filter behaviour for users who did not always know the right administrative terminology – and who needed to trust the results – maps directly onto the archivist workflow problem Preservica is solving.

Earlier in my career I implemented document management systems both at my own consultancy and for CLINET Platforms, and I worked as Proxy Product Owner on the migration of a complex internal catalogue at AOK health insurance. These were not glamorous projects, but they gave me a concrete understanding of how document collections degrade over time, how metadata becomes inconsistent across organisational handovers, and how users develop workarounds when search fails them. That context makes me a more useful partner to the

engineering team when trade-offs between indexing strategy, relevance ranking, and UI design need to be made with real user behaviour in mind.

I work AI-native by default: I use Claude and GPT for specification drafting and edge-case generation, Lovable and Framer for rapid prototyping before engineering commitment, and I have built my own agentic workflow tooling in n8n. I understand AI capabilities and limitations at a level that lets me make sensible product decisions about AI-powered features – including the contexts where they introduce risk or need careful design guardrails. Shipping AI features that archivists and records managers can actually trust is a harder problem than shipping AI features that impress in a sales demo, and I take that distinction seriously.

I am based in Berlin, authorised to work in the UK remotely as an EU citizen, and available within four weeks of an offer.

Best regards, *Philipp Hilbert*
