

Philipp Hilbert. *Platform PM who ships self-serve and makes complex products feel obvious.*

Positioning, ten years owning B2B SaaS platforms end-to-end, from admin config to integrations to agentic workflows, now applied to the console that puts 1mind's Superhumans in the hands of customers.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear 1mind team,

The Console role caught my attention because it is rare: a product brief that is simultaneously about information architecture, integration plumbing, permissions models, and an AI assistant that takes action on a customer's behalf, all in one scope. That is not a role for someone who wants to sit in one lane. I have spent the last decade moving between those lanes, and I would like to explain why I think that background maps onto what you need.

At CLINET Platforms I took over a three-person team facing a roadmap it could not deliver. I cut it to the top ten percent of features, user acceptance rose more than fivefold. The discipline there was not ruthlessness; it was constant re-anchoring to the real problem rather than the stated request. The equivalent at 1mind is the brief to "separate one-off requests from the broader product problems worth solving" in partnership with Customer Success. I know how that conversation goes, and I know how to run it without making CS feel bypassed.

At EMIL Group I joined an 18-person delivery organisation that was haemorrhaging senior engineers and client trust simultaneously. I restructured it into three lean teams with their own leads, attrition collapsed to a single departure, and I shipped the Claims Center from concept to production, cutting time-to-market for a new insurance product from many months to under a week. The structural work and the shipping work were not separate phases; they ran in parallel because the team could not afford to pause. I expect the same at 1mind: the Console needs to become more intuitive while the company is still growing fast, not after it has stabilised.

At Bundesagentur für Arbeit I owned a platform that spanned a Jenkins CI pipeline, ETL orchestration with Dagster and dbt, a UI relaunch, and a container migration, all concurrently, while serving as Deputy IT Security Officer. The integration-ownership and permissions-model thinking the Console role calls for is not new territory for me. I am comfortable sitting at the boundary between the UI a customer sees and the configuration model that makes it possible, and I can hold that conversation with an engineer in a way that produces a durable solution rather than a one-screen fix.

On the AI assistant dimension: I have built multi-agent frameworks and agentic workflow tooling for my own projects, and I use AI prototyping tools, Claude Code, Lovable, Cursor, as my daily working environment, not as an experiment. I understand the trust and transparency problems that arise when an assistant takes action on a user's behalf, and I have thought carefully about how to make those actions legible. That is the right starting point for defining what the 1mind Console assistant can and cannot do.

I am based in Berlin and available to work fully remotely in the Eastern timezone overlap, which I do routinely with US-based counterparts.

Best regards, *Philipp Hilbert*
