

# Philipp Hilbert. *Execution-first Product Manager for AI-driven enterprise platforms.*

*Positioning*, ten years translating complex operational workflows into engineering-ready software, now applying that to procurement intelligence and AI-native enterprise tooling.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · [www.true-north.berlin](https://www.true-north.berlin) · [YouTube \(EN sample\)](#)

---

Dear Green Cabbage team,

Procurement workflows carry enormous hidden complexity: approval chains, supplier data quality, contract version drift, and the gap between what buyers know and what the data actually says. I have spent the last several years working in analogous environments, specifically AI data management for European defence systems at Rohde & Schwarz and public labour-market analytics at the Federal Employment Agency, where the core challenge was the same as yours: take messy, high-stakes operational data and build software systems that make it actionable for non-technical users who need to trust what they see.

At Rohde & Schwarz I owned a central AI data platform as the sole Product Manager on a government contract, delivering end-to-end from user stories through GitOps rollout. I reduced deployment times from multiple days to under forty minutes by introducing ArgoCD-based automation, and I unblocked an infrastructure issue that had been outstanding for over a year by stepping in technically and delivering the feature myself in one week. That combination of tight requirements discipline and enough technical depth to remove blockers without waiting is precisely what your job description describes when it asks for someone who can "clarify requirements, refine scope, and remove ambiguity" alongside engineering teams.

At the Federal Employment Agency I owned an analytics platform serving regional labour market data to government and enterprise users. I introduced Dagster for automated ETL and indicator generation, designed interactive data dashboards and comparison tooling, and served as Deputy IT Security Officer with responsibility for compliance and data lineage. Building governed, auditable data pipelines for enterprise users with high trust requirements translates directly to procurement intelligence, where the credibility of the underlying data is a core part of the product value proposition.

On the AI side, I have worked hands-on with agentic workflow design using n8n and file-queue-based agent loops, LLM-assisted specification and edge-case generation using Claude and GPT, and rapid front-end prototyping in Lovable and Framer. I do not pitch AI features for novelty. My standard at Rohde & Schwarz was that every AI capability had to be grounded in measurable operational value and auditable data provenance before it reached a stakeholder demo, which matches Green Cabbage's stated standard of AI features being "grounded in measurable operational value and enterprise usability."

I am looking for a role where requirements discipline and AI-native execution are both valued, and where the product has real operational stakes for its users. Green Cabbage's Harvest platform is exactly that. I would welcome the opportunity to discuss the role in more detail.

Best regards, *Philipp Hilbert*

---