

Philipp Hilbert. *Platform PM who speaks both FHIR and Kubernetes.*

Positioning, ten years building data platforms in regulated environments – now bringing that depth to healthcare interoperability at the infrastructure layer.

BERLIN, GERMANY · HILBERT@TRUE-NORTH.BERLIN · TRUE-NORTH.BERLIN

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

§ WHY 1UPHEALTH

1upHealth sits at the exact intersection where my experience is densest: cloud-native data platforms, regulated data flows, and the engineering complexity of ingesting and normalising data at scale. The posting calls out FHIR interoperability and ETL pipeline ownership as core to the role – those are precisely the patterns I worked with at the Federal Employment Agency building the Arbeitsmarktmonitor's Dagster-based ETL layer and at Rohde & Schwarz designing reproducible data transformation pipelines with full audit trails. A platform that routes clinical and claims data across payers, providers, and digital health companies is a harder version of the same coordination problem I have been solving across regulated industries for the past several years. *That is exactly the environment where my strengths compound.*

§ HONEST FRAMING

I have not shipped a product inside a health system using FHIR or HL7v2 directly; my healthcare-adjacent work – the CLINET mobile hospital app and the AOK ITSCare proxy role – touched clinical workflows without owning the interoperability layer. The transferable mechanic I bring is strong: designing ETL pipelines, data provenance systems, and API-centric workflows in comparably regulated, compliance-heavy environments, and translating that architecture for both technical and non-technical stakeholders.

§ AI-NATIVE STACK

Daily tools, not buzzwords.

- **Claude**, spec drafting, edge-case generation, async reasoning
- **GPT**, research synthesis and user-story iteration
- **GitHub Copilot**, in-editor accelerated implementation
- **Claude Code**, in-repo refactors and acceptance-criteria validation
- **Cursor / VSCode**, day-to-day development environment
- **Lovable**, rapid front-end prototyping
- **Framer**, design-heavy wireframing and stakeholder demos
- **Gamma.app**, roadmap and narrative decks

§ TECHNICAL ENVIRONMENT

Shipped with, not just listed.

Python, Kotlin (Spring Boot), JavaScript, PostgreSQL, MongoDB, REST APIs, Dagster, dbt, DuckDB, Redpanda, Kubernetes, ArgoCD, Argo Workflows, Helm, OpenShift, Jenkins, Docker, GitOps, MicroStrategy, AngularJS, Ionic, CGM API integration.

§ SIDE PROJECT

BACKTESTING ENGINE [GITHUB.COM/HILBERTP](https://github.com/HilbertP) →

Self-built crypto backtesting and agent framework for systematic strategy validation.

- Ingests historical market data, executes strategy logic, and produces performance metrics across configurable time windows.
- Designed around a file-queue-based agent loop architecture for parallelisable, reproducible runs.
- Active development with real trade data from Hyperliquid and Bitget feeding validation cycles.
- Demonstrates the same pipeline-design and data-provenance thinking the 1upHealth platform requires — just applied to market microstructure rather than clinical data.

§ HOW I WOULD WORK AT 1UPHEALTH

First weeks, concrete.

WEEK 1

Listen, map, find the gaps.

Shadow customer success calls and engineering standups to map where data ingestion breaks down, where FHIR normalization creates ambiguity, and where the existing roadmap has unresolved dependencies between the API layer and downstream analytics consumers.

WEEK 2

Slice, validate, prioritise.

Consolidate findings into a prioritised gap map, validate the top three with an engineer and one customer-success counterpart, and surface a concrete recommendation on sequencing to the team.

WEEK 3+

Ship, close loops.

Own the first well-scoped user story through refinement, acceptance criteria, and delivery – establishing a working rhythm with engineering before touching the longer-horizon roadmap.

I COVER LETTER

To the 1upHealth Product team,

Healthcare data is one of the few domains where poor data quality directly harms people rather than just degrading conversion rates. That framing matters to me, and it is part of why the PM role at 1upHealth is compelling. You are building the infrastructure layer that lets clinical and claims data move correctly across an ecosystem that has historically treated interoperability as an afterthought. That is a hard, important problem, and it requires a PM who can hold the technical detail and the customer story at the same time.

My most directly transferable experience is at Rohde & Schwarz, where I owned the AI and data management platform for the FCAS defence programme. The core challenge there was the same one 1upHealth faces: ingesting multimodal data from heterogeneous sources, normalising it into a consistent structure, and making it auditable enough that downstream ML workflows could trust the provenance of every record. I designed the transformation pipelines, wrote the data provenance and audit-trail requirements, and introduced GitOps-based rollout that cut

deployment times from several days to under forty minutes. I did that as effectively the sole PM on a cross-functional team of architects, data scientists, and backend engineers – which meant constant translation between what compliance and stakeholders needed and what engineering could feasibly build in a given sprint.

At the Federal Employment Agency I added ETL pipeline ownership in a more analytics-heavy context: introducing Dagster for automated indicator generation, migrating the Arbeitsmarktmonitor to Kubernetes, and serving as Deputy IT Security Officer. The combination of data-lineage accountability and delivery ownership is directly relevant to a platform that processes sensitive clinical records for payers and providers operating under strict regulatory obligations.

Earlier, at Emil Group, I took on a struggling InsurTech SaaS platform mid-delivery. I rebuilt sprint discipline, split an unwieldy seventeen-person team into effective Scrum units, re-established velocity tracking, and ran weekly product workshops to rebuild client trust. The Claims Center module I designed from concept to production rollout in that environment is the kind of end-to-end ownership the 1upHealth PM role describes. I mention it because the ability to stabilise delivery while simultaneously advancing the product roadmap is a rare combination, and it is one I have demonstrated under pressure.

I am honest about where my gap sits: I have not yet shipped a product with FHIR or HL7v2 as the primary integration layer. What I bring instead is a track record of learning regulated technical standards quickly, designing around their constraints, and making them legible to both clinical and engineering audiences. Given that 1upHealth explicitly values startup pragmatism and the ability to translate complex technical topics for mixed audiences, I believe the transferable mechanics are strong.

I would be glad to walk through any of these projects in more detail, or to discuss how I would approach the first roadmap conversation with your engineering and customer success teams.

Best regards, *Philipp Hilbert*

II CURRICULUM VITAE

§ BEARING

Product Manager with nearly ten years of experience owning data platforms, ETL pipelines, and API-centric systems in regulated environments. Background spans defence AI data infrastructure, public-sector analytics platforms, InsurTech SaaS, and mobile healthcare applications. Comfortable translating complex technical architectures for clinical, business, and

engineering audiences. Fluent in AI-native workflows, backlog discipline, and cross-functional delivery. Brings founder-level pragmatism and a consistent track record of unblocking teams and shipping end-to-end in ambiguous, compliance-sensitive contexts.

§ CORE STRENGTHS

- ETL pipeline and data provenance ownership in regulated systems
 - Translating technical data flows into business-legible requirements
 - API-centric platform PM: REST, integration design, consumer personas
 - Cross-functional delivery in agile environments under pressure
 - Data quality, audit trails, and lineage accountability
 - Roadmap sequencing with explicit trade-off reasoning
 - Stakeholder alignment across engineering, compliance, and customer success
 - AI-native execution: spec-to-delivery with LLM-assisted tooling
 - Founder pragmatism: high agency, low ego, full delivery ownership
-

§ PROFESSIONAL EXPERIENCE

ROHDE & SCHWARZ 2024, 2025

Product Manager, AI Data Transformation

- Designed reproducible data transformation pipelines with full audit trails and data provenance for a multimodal AI training platform (FCAS defence programme)
- Enabled ML training and evaluation on curated datasets with transparent evaluation concepts (SafeAI, ExplainableAI)
- Delivered autonomous feature sets end-to-end: user stories, acceptance criteria, implementation support, testing, and GitOps-based rollout
- Reduced deployment times from several days to 20–40 minutes by introducing GitOps with ArgoCD
- Led cross-functional team: architect, data scientists, backend and frontend engineers, QA, multinational stakeholders

FEDERAL EMPLOYMENT AGENCY 2023, 2024

Product Manager, Public Data Analytics

- Introduced Dagster for automated ETL and indicator generation on the Arbeitsmarktmonitor analytics platform

- Delivered labour shortage analytics, flexible regional comparison tools, and modernised admin tooling
- Served as Deputy IT Security Officer with focus on compliance, governance, and data lineage
- Migrated legacy systems to Kubernetes; standardised CI/CD with full test integration at build time
- Designed interactive dashboards and community features for regional labour market data consumers

EMIL GROUP 2022, 2022

Product Manager, InsurTech SaaS Platform

- Designed and delivered a Claims Center module from concept to production rollout on a B2B InsurTech platform serving primary insurers, reinsurers, MGAs, and brokers
- Delivered modules across product configuration, pricing, underwriting, policy issuance, document management, and workflow automation
- Rebuilt sprint discipline: split a 17-person team into focused Scrum units, reintroduced velocity tracking and structured refinement
- Regained client trust in a prototype-stage environment through structured weekly product workshops
- Reduced senior developer attrition from near-total to a single departure

CLINET PLATFORMS 2021, 2022

Product Owner, Mobile Healthcare App

- Delivered core flows for digital anamnesis, meal and therapy plans, transport, chat, and patient document storage on a mobile-first hospital pilot app
- Integrated CGM interface into hospital information system workflows, managing legacy API constraints transparently with stakeholders
- Led small team under strict deadline pressure with explicit prioritisation on highest-leverage features

§ FOUNDER & EARLY-STAGE EXPERIENCE

QCRYPT AG 2016, 2018

Product Owner, Co-Founder

- Owned backup, restore/checkpoint, and data-encryption product architecture for a quantum-secure encryption product based on one-time-pad keys and true random number generator hardware

- Served as deputy IT security officer; designed attack-vector mitigation into the product architecture
- Setup experience under two minutes with no technical knowledge required – drove the usability constraint as a hard product requirement

KVITT PAYMENT SOLUTIONS 2013, 2018

Founder, CFO

- Designed full UX/UI and mobile app prototype for group-based P2P payments on a P2P backbone
- Owned PSP integration and payment edge-case handling; product later acquired and operated by Sparkasse

BUNDESDRUCKEREI – COVID-19 DATA PLATFORM 2021, 2021

Product Owner

- Delivered secure data structures and identity management for RKI projects DIM and DESH, operating under strict federal data governance requirements

§ LANGUAGES

German, native. English, C2.

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

TRUE NORTH · APPLICATION PACKAGE · 1UPHEALTH

