

Philipp Hilbert. *Data platform PM who ships governed AI products.*

Positioning, ten years building data-intensive B2B platforms – now focused on governed AI capabilities inside enterprise data management.

BERLIN, GERMANY · HILBERT@TRUE-NORTH.BERLIN · TRUE-NORTH.BERLIN

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

§ WHY SEMARCHY

Semarchy sits at a genuinely interesting inflection point: MDM and data governance platforms are being asked to become the trusted foundation for enterprise AI, and the teams that figure out agentic, MCP-based access to governed data products first will define the category for the next decade. The job description names AI-assisted stewardship and agentic access to governed Data Products as the near-term focus – that is precisely the intersection where my experience in AI data transformation (Rohde & Schwarz, FCAS programme) and data platform ownership (Bundesagentur für Arbeit) lands. Working across French and English-speaking colleagues in a remote-first, genuinely international team is a setup I have operated in before and prefer. *That combination of governed data infrastructure and frontier AI integration is exactly the problem space where my strengths compound.*

§ HONEST FRAMING

I have not held a role explicitly titled MDM or data catalog PM, so I come to master data management from the adjacent disciplines of AI data provenance, ETL platform ownership, and audit-trail design rather than from a pure MDM background. The transferable mechanic is strong: designing governed data pipelines with full lineage, curating datasets for ML training with ExplainableAI constraints, and translating complex data-quality requirements into backlog-ready specifications are the same cognitive moves MDM product ownership demands.

§ AI-NATIVE STACK

Daily tools, not buzzwords.

- **Claude**, discovery synthesis and specification drafting
- **GPT**, market research and edge-case generation
- **GitHub Copilot**, in-repo implementation support
- **Cursor / VSCode**, LLM-assisted code review and feature scoping
- **Lovable**, rapid front-end prototyping
- **Framer**, design-heavy concept validation
- **n8n**, multi-agent orchestration and workflow automation
- **Gamma.app**, stakeholder-facing roadmap and briefing decks

§ TECHNICAL ENVIRONMENT

Shipped with, not just listed.

Python, Kotlin (Spring Boot), JavaScript, PostgreSQL, MongoDB, REST APIs, dbt, Dagster, DuckDB, Kubernetes, ArgoCD, Argo Workflows, Helm, OpenShift, Docker, Jenkins, GitOps pipelines, Redpanda, MicroStrategy, CGM API integration, Shopware, AngularJS, Ionic, low-code insurance platform tooling.

§ SIDE PROJECT

CRYPTO BACKTEST ENGINE [GITHUB.COM/HILBERTP](https://github.com/hilbertp) →

Personal infrastructure project: an algorithmic backtesting and agent framework for crypto markets.

- Ingests raw market microstructure data (order books, liquidation maps, funding rates) and runs configurable strategy simulations.
- Built on a file-queue-based agent loop architecture with modular strategy plugins, enabling parallel backtest runs without a heavyweight orchestration dependency.
- Reflects hands-on familiarity with execution concepts – slippage, price impact, routing – applied to perpetual DEX and spot market data.
- Demonstrates the same data pipeline instincts (provenance, reproducibility, auditability) that governed data product work requires.

§ HOW I WOULD WORK AT SEMARCHY

First weeks, concrete.

WEEK 1

Listen, map, find the gaps.

Shadow customer calls and field conversations, map the current discovery backlog against the job description's stated focus on agentic and AI-assisted stewardship, and identify where customer pain intensity has not yet been formally validated.

WEEK 2

Validate, slice, frame.

Run structured problem-discovery sessions with two or three customers already engaged in design-partner programmes, validate the highest-friction points in governed Data Product access, and produce a first problem-framing document for the engineering lead.

WEEK 3+

Ship, enable, close.

Produce the first shippable specification for one AI-oriented capability, align with engineering on acceptance criteria, and draft the internal enablement brief so the field can speak to it before it lands.

I COVER LETTER

To the Semarchy Product team,

Data platforms are being asked to do something they were not originally designed for: serve as the trusted, governed foundation for enterprise AI systems that act autonomously. That shift – from passive data repository to active participant in agentic workflows – is precisely the product problem I find most interesting right now, and it is the reason the Semarchy Product Manager role caught my attention immediately.

My most directly relevant work was at Rohde & Schwarz, where I owned the product scope for a central AI and data management platform inside the FCAS European defence programme. The core challenge was identical in structure to what Semarchy faces: multimodal datasets needed to be discoverable, trustworthy, and auditable before ML training could happen on them. I designed and shipped reproducible data transformation pipelines with full provenance, introduced SafeAI and ExplainableAI evaluation concepts, and delivered the end-to-end feature

set – from user stories through GitOps rollout – largely autonomously. Before that, at the Federal Employment Agency, I owned an analytics data platform, introduced Dagster for automated ETL, and served as Deputy IT Security Officer with explicit responsibility for compliance, governance, and data lineage. Both roles required exactly the discipline the Semarchy job description names: anchoring decisions in customer outcomes, producing specifications that engineers can act on, and maintaining governance without sacrificing velocity.

On the AI-native and agentic side: I have designed and operated multi-agent orchestration workflows using n8n and file-queue-based agent loops, built with LLM-assisted development tooling daily, and prototyped features rapidly in Lovable and Framer before committing engineering time. I am genuinely engaged with MCP as an emerging standard for agentic data access – not as a buzzword but as a concrete integration pattern I have thought through in the context of governed enterprise data. I can have a real technical conversation with engineers about architecture tradeoffs, and I can turn the same conversation into a customer-facing brief the same afternoon.

On the commercial side: at Emil Group I ran the full PM lifecycle for a B2B InsurTech SaaS platform, including re-establishing sprint discipline for a 17-person team that had lost delivery rhythm, rebuilding client trust through structured workshops, and shipping a Claims Center module from concept to production. That experience of holding the customer relationship, the engineering relationship, and the roadmap simultaneously – under real delivery pressure – is the kind of preparation that makes field enablement and commercialisation retrospectives feel like a natural part of the role rather than a burden.

I am fluent in English and German; French is not yet at the level I would claim for customer-facing work, though I can read product documentation and follow technical discussions in French and I am investing in improving. I wanted to name that gap directly rather than let it surface later.

Semarchy is at the moment where governed data platforms and enterprise AI need to be designed together, not bolted together after the fact. That is a product problem worth solving carefully, and I would like to help solve it.

Best regards, *Philipp Hilbert*

II CURRICULUM VITAE

§ BEARING

Product Manager with nearly ten years of experience across AI data platforms, analytics infrastructure, B2B SaaS, and regulated environments. Most recent work centred on governed data pipelines, ML dataset curation, and audit-trail design inside a major European defence AI programme, plus data platform ownership and IT security governance at a federal public institution. Comfortable owning product scope end-to-end – discovery, specification, delivery, and field enablement – in environments where requirements are still forming. AI-native operator: LLM-assisted development, rapid prototyping, and multi-agent workflow design are part of everyday practice. Low ego, high agency, founder-level pragmatism applied inside larger organisations.

§ CORE STRENGTHS

- Customer discovery and pain-intensity validation in B2B enterprise contexts
 - Governed data pipeline design with provenance, lineage, and auditability
 - AI-native product execution and agentic workflow design
 - End-to-end specification production: user stories through acceptance criteria
 - Roadmap ownership with transparent trade-off communication
 - Cross-functional alignment across engineering, data science, QA, and compliance
 - Sprint discipline, backlog management, and delivery stability under pressure
 - Technical fluency: databases, APIs, cloud infrastructure, ETL systems
 - Field enablement, stakeholder workshops, and customer-facing documentation
-

§ PROFESSIONAL EXPERIENCE

ROHDE & SCHWARZ 2024, 2025

Product Manager, AI Data Transformation

- Designed and shipped reproducible data transformation pipelines with full provenance and audit trails for multimodal datasets (image, audio, tabular) within the FCAS European defence AI programme
- Delivered SafeAI and ExplainableAI evaluation concepts enabling governed ML training on curated datasets
- Delivered autonomous feature sets end-to-end: user stories, acceptance criteria, implementation support, testing, and GitOps-based rollout

- Reduced deployment times from several days to 20–40 minutes by introducing GitOps with ArgoCD
- Led cross-functional team including architect, data scientists, backend and frontend engineers, and QA with multinational stakeholders

Tech: Kotlin (Spring Boot), Python, PostgreSQL, MongoDB, Redpanda, Kubernetes, Argo Workflows, ArgoCD, OpenShift

FEDERAL EMPLOYMENT AGENCY (BUNDESAGENTUR FÜR ARBEIT) 2023, 2024

Product Manager, Public Data Analytics

- Owned the Arbeitsmarktmonitor analytics platform for regional labour market and demographic data, delivering labour shortage analytics, regional comparison features, and modernised admin tooling
- Introduced Dagster for automated ETL and indicator generation; designed interactive dashboards and community features
- Served as Deputy IT Security Officer with explicit responsibility for compliance, governance, and data lineage
- Migrated legacy systems to Kubernetes and standardised CI/CD with full test integration at build time

Tech: Python, Django, JavaScript, PostgreSQL, dbt, Dagster, DuckDB, MicroStrategy, Kubernetes, Jenkins

EMIL GROUP 2022, 2022

Product Manager, InsurTech SaaS Platform

- Owned product scope across configuration, pricing, underwriting, policy issuance, document management, and workflow automation for a B2B InsurTech platform serving primary insurers, reinsurers, MGAs, and brokers
- Designed and delivered a Claims Center module from concept to production rollout
- Re-established sprint discipline and delivery stability; split a bloated 17-person team into effective Scrum units and introduced structured ownership and deputy concepts
- Regained client trust in a prototype-stage environment through structured communication and weekly product workshops

Tech: Low-code platform, Java, HTML/CSS, Kubernetes, REST APIs

CLINET PLATFORMS 2021, 2022

Product Owner, Mobile Healthcare App

- Delivered core flows for digital anamnesis, meal and therapy plans, transport, chat, and patient document management on iOS and Android
- Integrated CGM interface into hospital information system workflows; managed legacy API and external database complexity transparently with stakeholders

Tech: AngularJS, Ionic, Python, Kubernetes, CGM API integration

§ FOUNDER & EARLY-STAGE EXPERIENCE

BUNDESDRUCKEREI 2021, 2021

Product Owner, COVID-19 Data Platform

- Owned product scope for RKI projects DIM and DESH: secure data structures, identity management, and data governance.

QCRYPT AG 2016, 2018

Co-Founder and Product Owner

- Defined and shipped a quantum-secure encryption product (one-time-pad keys, true random number generator hardware) with a sub-two-minute setup UX requiring no technical knowledge from end users.
- Served as deputy IT security officer; responsible for attack-vector mitigation and security architecture review.

KVITT PAYMENT SOLUTIONS 2013, 2018

Founder and CFO

- Built a group-based P2P payment product on a P2P backbone: PSP integration, full UX/UI design, mobile app prototype. Later acquired and operated by Sparkasse.
-

§ LANGUAGES

German, native. English, C2.

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

