

Philipp Hilbert. *Data platform PM who ships governed AI products.*

Positioning, ten years building data-intensive B2B platforms – now focused on governed AI capabilities inside enterprise data management.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

To the Semarchy Product team,

Data platforms are being asked to do something they were not originally designed for: serve as the trusted, governed foundation for enterprise AI systems that act autonomously. That shift – from passive data repository to active participant in agentic workflows – is precisely the product problem I find most interesting right now, and it is the reason the Semarchy Product Manager role caught my attention immediately.

My most directly relevant work was at Rohde & Schwarz, where I owned the product scope for a central AI and data management platform inside the FCAS European defence programme. The core challenge was identical in structure to what Semarchy faces: multimodal datasets needed to be discoverable, trustworthy, and auditable before ML training could happen on them. I designed and shipped reproducible data transformation pipelines with full provenance, introduced SafeAI and ExplainableAI evaluation concepts, and delivered the end-to-end feature set – from user stories through GitOps rollout – largely autonomously. Before that, at the Federal Employment Agency, I owned an analytics data platform, introduced Dagster for automated ETL, and served as Deputy IT Security Officer with explicit responsibility for compliance, governance, and data lineage. Both roles required exactly the discipline the Semarchy job description names: anchoring decisions in customer outcomes, producing specifications that engineers can act on, and maintaining governance without sacrificing velocity.

On the AI-native and agentic side: I have designed and operated multi-agent orchestration workflows using n8n and file-queue-based agent loops, built with LLM-assisted development tooling daily, and prototyped features rapidly in Lovable and Framer before committing engineering time. I am genuinely engaged with MCP as an emerging standard for agentic data access – not as a buzzword but as a concrete integration pattern I have thought through in the context of governed enterprise data. I can have a real technical conversation with engineers about architecture tradeoffs, and I can turn the same conversation into a customer-facing brief the same afternoon.

On the commercial side: at Emil Group I ran the full PM lifecycle for a B2B InsurTech SaaS platform, including re-establishing sprint discipline for a 17-person team that had lost delivery rhythm, rebuilding client trust through structured workshops, and shipping a Claims Center module from concept to production. That experience of holding the customer relationship, the engineering relationship, and the roadmap simultaneously – under real delivery pressure – is the kind of preparation that makes field enablement and commercialisation retrospectives feel like a natural part of the role rather than a burden.

I am fluent in English and German; French is not yet at the level I would claim for customer-facing work, though I can read product documentation and follow technical discussions in French and I am investing in improving. I wanted to name that gap directly rather than let it surface later.

Semarchy is at the moment where governed data platforms and enterprise AI need to be designed together, not bolted together after the fact. That is a product problem worth solving carefully, and I would like to help solve it.

Best regards, *Philipp Hilbert*
