

Philipp Hilbert. *Product Manager, internal tools and automation.*

Positioning, ten years shipping workflow automation and data platforms for demanding internal users, now bringing that stack to CapsLock's media buying engine.

BERLIN, GERMANY · HILBERT@TRUE-NORTH.BERLIN · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear CapsLock team,

I have spent the last decade owning products where the user is a domain expert with a tight feedback loop and the outcome lands in a number you can read the same day, pipeline run counts, data update time, time-to-market for an insurance product. The Media Buyer Automation Platform is that kind of product, and the description is specific enough that I can tell the team has already mapped the problem clearly: pacing logic, spend guardrails, anomaly alerts, approval flows, bulk actions. That specificity is a good sign. It means I would be walking into a context where the hard discovery has been done and the job is to own delivery, drive adoption, and keep the roadmap honest.

At Rohde & Schwarz I inherited an AI and data platform where rollout took days and the UI that had been built with great care recorded zero production clicks. I instrumented both workflows, redirected the roadmap toward the developer-code-first path users actually used, and brought rollout time down to 30 minutes via ArgoCD GitOps. The lesson, that usage data beats opinion every time, and that a beautiful feature nobody touches is a failed feature, is directly applicable to spend-affecting automation logic where a wrong guardrail is worse than no guardrail.

At the Bundesagentur für Arbeit I owned a national labour-market analytics platform and introduced Dagster and dbt so quarterly data imports ran without hand-holding, cutting data update time by roughly 75%. The parallel to CapsLock's reporting and analytics workflows is direct: the goal in both cases is to get reliable, timely data in front of decision-makers without requiring an engineer on standby.

Beyond the core role, the posting mentions active use of AI tools with examples requested. My daily stack includes Claude Code for in-repo spec work and edge-case generation, Cursor for code navigation, Lovable for rapid UI prototyping, n8n for multi-step workflow automation, and GitHub Copilot for test generation. I also built construct8, a two-sided B2B marketplace, solo, product, automation, and funnel analysis, which is the closest thing I have to a portfolio piece for the builder orientation the role describes.

I am comfortable with the scope evolving. The framing around portfolio growth across marketing automation, reporting, creative operations, and cross-platform account management matches how I have worked before: one core platform as an anchor, adjacent products added as capacity and trust are established. I would welcome a conversation about the current roadmap priorities and how the team makes sequencing decisions.

Best regards, *Philipp Hilbert*
