

Philipp Hilbert. *Open-source observability platform product leader.*

Positioning, ten years shipping infrastructure and data platform products for technical users, now bringing that depth to the open-source database monitoring space.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

§ WHY PERCONA

PMM sits at an interesting intersection: a beloved open-source tool with real enterprise gravity, used by exactly the kind of technical users, DBAs, SREs, platform engineers, who will tell you the truth about what is broken and why. The posting's explicit call for someone who talks to customers early, runs cheap experiments, and ships before the spec is perfect matches how I have worked at every stage of my career. The detail that PMM's users rely on it for mission-critical database environments means the discovery bar is high and the feedback loop is tight.

That is exactly the environment where my strengths compound.

§ HONEST FRAMING

PMM's primary observability layer is built around Prometheus, Grafana, and VictoriaMetrics, tools I have read about and tracked closely but have not shipped production features against. My observability hands-on time has lived one layer up: instrumenting CI/CD pipelines, owning ETL and data platform products, and driving metric-based product decisions. The underlying skill, translating what monitoring data means for a technical user into a prioritised roadmap, is one I have exercised continuously.

§ AI-NATIVE STACK

Daily tools, not buzzwords.

- **Claude Code**, in-repo refactors and spec generation
- **GitHub Copilot**, inline code assistance across Python and SQL
- **Cursor / VSCode**, environment for agent-assisted development
- **Lovable**, rapid front-end prototyping and stakeholder demos
- **n8n**, agentic workflow orchestration and automation
- **GPT**, research, edge-case generation, and document drafting
- **Gamma.app**, internal decks and product narrative
- **Framer**, design-heavy landing pages and prototype surfaces

§ TECHNICAL ENVIRONMENT

Shipped with, not just listed.

Kubernetes, OpenShift, Docker, ArgoCD, GitOps rollout and rollback, Jenkins, GitLab, Azure DevOps, CI/CD pipelines, Python, Kotlin (Spring Boot), Django, PostgreSQL, MongoDB, REST APIs, Kafka, Dagster, dbt, SQL, ETL pipelines, Jira, Confluence, Figma, Miro.

§ SIDE PROJECT

CONSTRUCT8 [GITHUB.COM/HILBERTP](https://github.com/HilbertP) →

Two-sided B2B marketplace matching construction firms with workers, built solo from zero.

- Anonymity-until-payment model that protects each side until commitment is confirmed.
- WhatsApp and Telegram as the primary product surface, keeping acquisition costs near zero.
- LLM agents across the funnel so one person runs the full day-to-day operation.
- Demonstrates end-to-end product ownership, discovery, architecture, delivery, in a lean founder setting directly comparable to Percona's "wear multiple hats" expectation.

§ HOW I WOULD WORK AT PERCONA

First weeks, concrete.

WEEK 1

Listen, map, find the gaps.

Sit in on customer support tickets, community forum threads, and one or two live customer calls to understand where PMM users feel friction, query performance diagnostics, alerting noise, or onboarding, before touching the roadmap.

WEEK 2

Instrument and validate.

Map the existing product telemetry against what I heard in week one, identify the three highest-signal gaps between assumed and actual usage, and bring a ranked shortlist to engineering.

WEEK 3+

Ship and iterate.

Take the top-ranked item through discovery to a shippable slice, write the acceptance criteria, and close the first loop with the community or a design partner before moving to the next problem.

I COVER LETTER

Dear Percona team,

I am writing to apply for the Product Manager role for Percona Monitoring and Management. PMM is a product I respect: open-source, trusted by DBAs and SREs for monitoring mission-critical database environments, and genuinely difficult to build well because its users are demanding, informed, and will not accept a tool that wastes their time.

My background is in infrastructure and data platform products for technical users. At Rohde & Schwarz I owned an internal AI and data platform for the FCAS European defence programme, leading a cross-functional team of fifteen engineers and data scientists through discovery, prioritisation, and GitOps delivery. A user workflow analysis I ran there revealed that a UI-led feature had recorded exactly zero clicks in production, I redirected the roadmap accordingly, and pipeline runs increased by 200%. At the Bundesagentur für Arbeit I owned a national labour-market analytics platform, migrated a bare-metal system fully into containers, introduced Dagster and dbt, and reduced the time per data update by almost 75%. Those are the kinds of problems, data pipelines, observability instrumentation, developer tooling, that sit close to PMM's core.

I work with PostgreSQL and MongoDB in product contexts: writing queries, reviewing schema decisions, and sitting close enough to the engineering discussion to translate what a database performance problem means for a user into a requirement that engineering can act on.

Kubernetes has been part of my technical environment for several years; I have owned products running on K8s and OpenShift, governed rollout and rollback via ArgoCD, and used containerisation as a lever for team autonomy. The open-source community dimension of the role is one I take seriously. At Qcrypt I built an enterprise product as a co-founder and maintained it as chair of the board; at construct8 I built a marketplace from zero alone. In both cases, the fastest feedback came from users who had strong opinions and were not shy about sharing them, which is exactly the community PMM already has.

What I bring that I think is less common in a PM profile: I write Python, I read SQL fluently, and I can sit in an engineering design discussion and contribute rather than observe. I build agentic workflows with n8n and Claude Code in my own projects. When the posting mentions "defining AI-powered workflows or AI-assisted observability experiences," that is not a stretch goal for me, it is something I am already doing at the workbench level. I am also comfortable with the ambiguity and pace the posting describes. Making a decision with incomplete information, running a cheap experiment to validate it, and iterating based on what users actually do rather than what they say they will do, that is the posture I bring to every role.

I would be glad to talk through how my experience maps to PMM's current roadmap challenges, and to hear more about where the product is headed. Thank you for considering my application.

Best regards, *Philipp Hilbert*

II CURRICULUM VITAE

§ BEARING

Senior Product Manager with 13 years of product ownership across infrastructure, data platform, AI, and regulated B2B SaaS products. I have owned products built on Kubernetes, PostgreSQL, and MongoDB, led cross-functional teams through full delivery cycles, and consistently grounded roadmap decisions in user workflow analysis and measurable outcomes. For an open-source observability platform like PMM, the combination of technical fluency, discovery discipline, and cross-functional delivery experience is the core of what I bring.

§ WHAT I CAN DO

Platform and Infrastructure Product Management: Kubernetes (K8s), OpenShift, Docker, ArgoCD, GitOps, CI/CD, Jenkins, containerised environments, release gating, same-day hotfix paths.

Data Platform and AI: ETL (extract, transform, load), Dagster, dbt (data build tool), PostgreSQL, MongoDB, Kafka, LLM (large language model) agents, agentic workflows, Explainable AI (XAI), EU AI Act, AI governance.

Product Ownership: requirements engineering, user stories, acceptance criteria, backlog refinement, sprint planning, velocity and burndown, UAT, Jira, Confluence, Figma, Miro.

Product Management: continuous discovery, product vision and roadmap, prioritisation, MVP and PoC definition, A/B testing, user research and interviews, product lifecycle management, go-to-market support.

Agile and Scaled Delivery: Scrum, Kanban, SAgile (Scaled Agile Framework), ART (agile release train), PI (program increment) planning.

Analytics and Data: funnel and churn analysis, KPI (key performance indicator) definition, Matomo, MicroStrategy, SQL.

AI Products and Tooling: Claude Code, GitHub Copilot, Cursor, Lovable, n8n, prompt and context engineering, evaluation loops, human-in-the-loop design.

Governance and Business: GDPR (General Data Protection Regulation), Deputy IT Security Officer, identity management, budget ownership, ROI (return on investment) reasoning.

§ PROFESSIONAL EXPERIENCE

ROHDE & SCHWARZ · PRODUCT MANAGER · DEFENCE AI AND DATA PLATFORM 2024-2025

Led a 15-person cross-functional team building the internal AI and data platform for the FCAS European programme; GitOps rollout via ArgoCD cut deployment time from days to 30 minutes, and a workflow-instrumentation exercise redirected the roadmap to increase pipeline runs by 200%.

BUNDESAGENTUR FÜR ARBEIT · PRODUCT MANAGER · PUBLIC-SECTOR ANALYTICS PLATFORM 2023-2024

Owned the national labour-market analytics platform; introduced Dagster and dbt to eliminate hand-holding on quarterly imports (saving ~75% of update time), migrated the system fully into containers, and carried a full UI relaunch and large ETL migration concurrently. Deputy IT Security Officer.

CLINET PLATFORMS · PRODUCT OWNER · HEALTHCARE MOBILE APP 2023

Cut an undeliverable roadmap to its top 10% against agreed success measures; user acceptance rose more than fivefold. Implemented the CGM interface into KIS hospital workflows, reducing digital anamnesis time from 30 minutes to 5.

EMIL GROUP · PRODUCT MANAGER · B2B SAAS INSURANCE PLATFORM 2022

Restructured an 18-person delivery organisation into three lean teams with their own leads, reducing attrition from near-total to a single departure. Shipped the Claims Center from concept to production and cut new-insurance-product time-to-market from months to under a week.

SCHEIDT & BACHMANN · PRODUCT OWNER · RETAIL PAYMENTS AND POS 2021-2022

Moved research results into production across a matrix organisation in a full SAFe setup with agile release trains and PI planning.

BUNDESDRUCKEREI · PRODUCT OWNER (INTERIM) · GOVERNMENT SECURE DATA PLATFORM 2021

Data structures and identity management for the Robert Koch Institute programmes DIM and DESH.

AOK ITSCARE · PRODUCT OWNER (PROXY) · HEALTH INSURANCE B2B PROCUREMENT 2019-2020

Migrated a legacy system to Shopware for three AOK insurers and ~20,000 employees; backlog refinement, user stories, and acceptance criteria in Jira.

ENERCON · PROJECT LEAD · WIND ENERGY SITE IT 2019

Full site IT modernisation end to end.

WALDBURG-ZEIL KLINIKEN · PRODUCT OWNER · CLINICAL NETWORK INFRASTRUCTURE 2017-2018

Wifi infrastructure rollout across nine hospitals.

§ FOUNDER & EARLY-STAGE EXPERIENCE

CONSTRUCT8 · FOUNDER AND PRODUCT MANAGER · CONSTRUCTION STAFFING MARKETPLACE 2025-2026

Zero to one alone: two-sided B2B marketplace with anonymity-until-payment model, WhatsApp/Telegram as product surface, and LLM agents across the funnel.

QCRYPT AG · CO-FOUNDER AND PRODUCT OWNER · QUANTUM-SECURE ENCRYPTION 2016-2021

Shipped a three-layer quantum-resistant encryption product that installs in under two minutes; later Chair of the board.

SMART SOIL TECHNOLOGIES · CO-FOUNDER, CDO AND CFO ·
NANOTECHNOLOGY AND AGRICULTURE 2018-2021

Took a fertiliser product from laboratory to investor rounds, owning financial planning and investor reporting.

KVITT PAYMENT SOLUTIONS · FOUNDER AND CFO · FINTECH P2P
PAYMENTS 2013-2018

Built from zero to acquisition; acquired and operated by Sparkasse inside the girocard ecosystem.

§ LANGUAGES

German, native. English, C2.

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

TRUE NORTH · APPLICATION PACKAGE · PERCONA

→