

Philipp Hilbert. *Open-source observability platform product leader.*

Positioning, ten years shipping infrastructure and data platform products for technical users, now bringing that depth to the open-source database monitoring space.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear Percona team,

I am writing to apply for the Product Manager role for Percona Monitoring and Management. PMM is a product I respect: open-source, trusted by DBAs and SREs for monitoring mission-critical database environments, and genuinely difficult to build well because its users are demanding, informed, and will not accept a tool that wastes their time.

My background is in infrastructure and data platform products for technical users. At Rohde & Schwarz I owned an internal AI and data platform for the FCAS European defence programme, leading a cross-functional team of fifteen engineers and data scientists through discovery, prioritisation, and GitOps delivery. A user workflow analysis I ran there revealed that a UI-led feature had recorded exactly zero clicks in production, I redirected the roadmap accordingly, and pipeline runs increased by 200%. At the Bundesagentur für Arbeit I owned a national labour-market analytics platform, migrated a bare-metal system fully into containers, introduced Dagster and dbt, and reduced the time per data update by almost 75%. Those are the kinds of problems, data pipelines, observability instrumentation, developer tooling, that sit close to PMM's core.

I work with PostgreSQL and MongoDB in product contexts: writing queries, reviewing schema decisions, and sitting close enough to the engineering discussion to translate what a database performance problem means for a user into a requirement that engineering can act on. Kubernetes has been part of my technical environment for several years; I have owned products running on K8s and OpenShift, governed rollout and rollback via ArgoCD, and used containerisation as a lever for team autonomy. The open-source community dimension of the role is one I take seriously. At Qcrypt I built an enterprise product as a co-founder and maintained it as chair of the board; at construct8 I built a marketplace from zero alone. In both cases, the fastest feedback came from users who had strong opinions and were not shy about sharing them, which is exactly the community PMM already has.

What I bring that I think is less common in a PM profile: I write Python, I read SQL fluently, and I can sit in an engineering design discussion and contribute rather than observe. I build agentic workflows with n8n and Claude Code in my own projects. When the posting mentions "defining AI-powered workflows or AI-assisted observability experiences," that is not a stretch goal for me, it is something I am already doing at the workbench level. I am also comfortable with the ambiguity and pace the posting describes. Making a decision with incomplete information, running a cheap experiment to validate it, and iterating based on what users actually do rather than what they say they will do, that is the posture I bring to every role.

I would be glad to talk through how my experience maps to PMM's current roadmap challenges, and to hear more about where the product is headed. Thank you for considering my application.

Best regards, *Philipp Hilbert*
