

Philipp Hilbert. *Technical Product Manager, AI Stockbroking.*

Positioning, 13 years shipping regulated financial products, from P2P payments to defence AI, now focused on building trustworthy AI-native stockbroking at scale.

BERLIN, GERMANY · HILBERT@TRUE-NORTH.BERLIN · TRUE-NORTH.BERLIN

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube](#) (EN sample)

Dear Bjak team,

I am applying for the Technical Product Manager, AI Stockbroking role. I have spent 13 years shipping regulated financial products, from a P2P payments product acquired into the girocard ecosystem to an AI and data platform for a European defence programme, and the thing that has stayed constant is the conviction that money-moving products have to be correct before they can be clever. Bjak's framing that reliability, regulatory requirements, and third-party financial rails are fundamental to the product, not considerations around the edges, is exactly the standard I hold myself to.

My most relevant direct preparation for this role sits at the intersection of fintech and AI evaluation. At Rohde & Schwarz I owned an internal AI and data platform for the FCAS programme: a cross-functional team of 15+ engineers and data scientists, GitOps delivery with ArgoCD, and a deliberate choice to instrument two competing design directions and kill the one that recorded zero actual usage. That discipline, measure the outcome, not the output, is the same discipline AI stockbroking features require, where a confident-sounding model answer that is subtly wrong is worse than no answer at all. I build evaluation loops before I build features.

On the regulated and payments side, I took Kvitt Payment Solutions from zero to acquisition by Sparkasse, operating inside the girocard ecosystem. I have shipped KYC and compliance workflows in insurance at EMIL Group, handled identity management for the Robert Koch Institute programmes at Bundesdruckerei, and worked through clinical data privacy requirements at CLINET Platforms. I understand that regulatory approvals are not bureaucratic overhead, they are launch gates, and the PM's job is to clear them in parallel with engineering, not in sequence. I have done that across insurance, health insurance, and public-sector analytics.

I am genuinely AI-native in practice, not in description. I use Claude Code and Copilot daily for in-repo work, I have built multi-agent orchestration workflows with n8n and file-queue-based loops, and I have prototyped front-end surfaces with Lovable and Framer fast enough to run real user tests before a sprint closes. More importantly, I know when AI should not be used: I would not put an LLM in control of a ledger entry without a deterministic guard and a human override path, and I can articulate that boundary clearly enough that engineering and compliance both trust the decision.

Bjak is expanding from insurance into stockbroking across Southeast Asia, a region where the infrastructure is real but the margin for error on user trust is thin. I want to be the PM who owns that vertical end to end, builds the evaluation and safeguard layer the AI features need to be trusted, and ships products where the user is never misled about the state of their funds. I move fast, I clear the gates that actually block launch, and I leave my ego at home. I would be glad to work through a practical assessment at your pace.

Best regards, *Philipp Hilbert*
