

Philipp Hilbert. *AI product leader for agentic data platforms.*

Positioning, Thirteen years shipping B2B products in regulated industries, now building with AI agents daily.

BERLIN, GERMANY · HILBERT@TRUE-NORTH.BERLIN · TRUE-NORTH.BERLIN

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear Matillion team,

Data engineering teams are bottlenecked, and the answer is an agent that can own the work, not just assist with it. That distinction matters: most AI tools in the data space are still in co-pilot mode, and the product challenge of getting to genuine autonomy, where a team actually trusts the output, is one I have been thinking about seriously.

My thirteen years of product ownership spans regulated and data-heavy environments: a national labour-market analytics platform at Bundesagentur für Arbeit, an AI and data platform for the FCAS European defence programme at Rohde & Schwarz, and a B2B SaaS platform for insurers at EMIL Group. In each case the hard work was not writing requirements, it was instrumenting workflows to find what users actually did versus what they said, then holding the quality bar under delivery pressure. At Rohde & Schwarz I ran two competing pipeline workflows in parallel and found the UI-led one recorded zero clicks. That kind of instrumented discovery is how I would approach Maia's eval and prioritisation work.

I use agentic workflows in my own product work every day, Claude Code, n8n orchestration, LLM agents across a live marketplace product I built solo. I am familiar with the operational tradeoffs in context management, tool use, and evaluation loops from building them, not just reading about them. When your posting mentions reasoning with Engineering about models, MCP, and orchestration, those are conversations I can contribute to, not just observe.

The data domain knowledge is a asset I can bring: ETL pipeline ownership with Dagster and dbt at Bundesagentur für Arbeit, container migration, and data platform product work across engineering and data science teams. The specific warehousing layer, Snowflake, Databricks, is territory I have not owned directly, but the architecture patterns and the user problems in data integration are familiar ground.

I would welcome the chance to talk about what the current quality gaps in Maia look like and how the team is measuring them.

Best regards, *Philipp Hilbert*
