

Philipp Hilbert. *AI coaching product from zero to shipped.*

Positioning. Thirteen years taking product from undefined to launched, now focused on AI-driven personalized experiences at consumer scale.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear Samsung Food team,

Samsung Health reaches hundreds of millions of people. The coaching capability you are building will decide whether that reach translates into moments that actually change behavior, or just noise that gets swiped away. That is a worthwhile problem to work on.

The job description names something I think about a lot: AI product failures are almost always context problems before they are model problems. What information reaches the model, when, in what format, and what is the minimum that produces a useful answer rather than the ideal one. I worked through that question on the FCAS AI and data platform at Rohde & Schwarz, where I redirected the product strategy after instrumenting two competing workflows and finding that the UI-led approach recorded exactly zero clicks. The developer-code-first path we switched to increased pipeline runs by 200%. The coaching equivalent is knowing when to send a message and when to stay quiet, and building the decision logic and evidence practice to tell the difference at scale.

I have taken products from undefined to launched across defence AI, public-sector analytics, insurance, and healthcare. At CLINET Platforms I cut a roadmap a five-person team could not deliver down to the top 10% of features, and user acceptance rose more than fivefold. At EMIL Group I shipped the Claims Center from concept to production with no precedent, while restructuring a delivery organisation that was losing its senior engineers. Both required the same thing: producing artefacts that survive being retold into other organisations without me in the room. That is what your job description calls "making decisions leave the room," and it is how I have worked across every multi-time-zone engagement I have had.

On the AI side, I build with these tools daily, Claude Code, GitHub Copilot, Cursor, Lovable, and I own the evaluation loop, not just the prompt. My construct8 project runs LLM agents across a full B2B marketplace funnel, with SQL-based funnel analysis to close the evidence loop. I understand the difference between a model that performs and a product that helps, and I know that the hard work is in the context design and the quality criteria, not the model selection.

Samsung Food's scale, tens of millions of users across phones, wearables, TVs, and appliances, and the depth of the Samsung Health data layer make this a rare opportunity to build coaching that could actually be useful at population scale. I would like to bring my product craft and my AI execution experience to that.

Best regards, *Philipp Hilbert*
