

Philipp Hilbert. *Product manager who ships AI-native EdTech at pace.*

Positioning, 13 years of end-to-end product ownership in B2B SaaS, with a hands-on AI-native stack and a record of turning complex workflows into products people actually use.

BERLIN, GERMANY · HILBERT@TRUE-NORTH.BERLIN · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Easygenerator's posting asks for a PM who is "excited by the potential but grounded in real customer problems." That phrase is the shortest description of how I have worked for the past thirteen years, so I want to show rather than claim it.

At Rohde & Schwarz I was brought in to accelerate an internal AI and data platform for the FCAS European defence programme. The team had built a genuinely elegant UI. When I got access to the usage instrumentation and ran it, the click count was exactly zero. I redirected the product to a developer-code-first design within weeks, not by overriding anyone, but by making the data impossible to ignore in the customer workshops. That is the AI-quality-and-adoption challenge in miniature: the thing that looks like the right AI experience is often not the one users actually reach for. Easygenerator's authoring tool lives inside the same tension, AI that feels powerful in a demo can feel intrusive or untrustworthy in a real content-creation workflow.

At CLINET Platforms I inherited a three-person team and a roadmap it could not deliver. I cut it to the top ten percent of features. User acceptance rose more than fivefold. At EMIL Group I restructured an 18-person delivery organisation that was losing its senior engineers: attrition fell from near-total exodus to a single departure. In both cases the lever was the same, ruthless prioritisation and a willingness to make the uncomfortable call before it becomes the expensive one. B2B SaaS EdTech at Easygenerator's stage needs exactly that discipline: a PM who protects the team's focus while the commercial organisation rightly wants everything at once.

On the AI side, I work hands-on: I build multi-agent orchestration frameworks, I prototype with Claude Code and Lovable before writing a spec, and I evaluate LLM outputs against real user acceptance criteria rather than vibes. I have designed human-in-the-loop checkpoints into agentic workflows, thought carefully about evaluation loops and trust signals, and read the EU AI Act for a client in a regulated environment. What I do not have is direct experience in the EdTech or learning-content domain. I expect a learning curve of a few weeks on the authoring-workflow specifics, and I am comfortable saying so.

Easygenerator is used by Kellogg's, T-Mobile, and Siemens, organisations where the person who holds the knowledge and the person who creates the training content are different people, and the gap between them is expensive. Closing that gap with AI that earns user trust rather than demanding it is a product problem I find genuinely compelling. I would be glad to talk through how I would approach the first ninety days in more detail.

Best regards, *Philipp Hilbert*
