

Philipp Hilbert. *AI Product Manager, investment intelligence at the frontier.*

Positioning, product manager with 13 years shipping regulated, data-intensive products who now builds AI-native systems end to end, from ETL pipelines to agent-based decision workflows.

BERLIN, GERMANY · HILBERT@TRUE-NORTH.BERLIN · TRUE-NORTH.BERLIN

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

§ WHY LIQUIDITY

Private credit is one of the few domains where deal quality and decision speed are genuinely in tension, and Liquidity is trying to dissolve that tension with proprietary decision science rather than just more analysts. The posting describes owning AI outputs across the full stack, data ingestion, model layers, agent-based origination and diligence workflows, which maps precisely to how I have operated on the FCAS AI and data platform and on the Bundesagentur analytics stack. The phrase "balance speed and precision in a financial context" is not a platitude here; in private credit, a mispriced deal or a missed window both cost real money, and the system has to handle both failure modes. *That is exactly the environment where careful AI product ownership compounds.*

§ HONEST FRAMING

Private credit as a domain, deal structuring, credit risk scoring, and lending-cycle workflows, is new territory for me; my financial product experience is in payments and insurance rather than capital markets. The transferable mechanic is strong: I have spent years translating complex, high-stakes decision processes (defence programme management, national labour-market policy) into production AI systems with strict accuracy and auditability requirements, which is the same underlying problem at a different asset class.

§ AI-NATIVE STACK

Daily tools, not buzzwords.

- **Claude Code**, in-repo refactors and agent scaffolding
- **GitHub Copilot**, pair-programming across Python and SQL
- **Cursor / VSCode**, context-aware editing and debugging
- **Lovable**, rapid front-end prototyping for discovery sprints
- **GPT**, spec generation and edge-case enumeration
- **n8n**, agentic workflow automation and multi-agent orchestration
- **Gamma.app**, stakeholder-ready decks from raw notes
- **Framer**, design-heavy prototypes for user testing

§ TECHNICAL ENVIRONMENT

Shipped with, not just listed.

Python, SQL, PostgreSQL, MongoDB, REST APIs, Kotlin (Spring Boot), Dagster, dbt, ETL pipeline design, Kubernetes, Docker, ArgoCD, GitOps, Jenkins, CI/CD, Azure DevOps, GitLab, Kafka, OpenShift, LLM agents, RAG, prompt engineering, evaluation loops, Explainable AI (XAI), EU AI Act compliance, Jira, Confluence, Miro.

§ SIDE PROJECT

CONSTRUCT8 [GITHUB.COM/HILBERTP](https://github.com/HilbertP) →

Two-sided B2B marketplace matching construction firms with workers, built solo from zero to one using LLM agents across the full funnel.

- Anonymity-until-payment model that protects both sides until commitment is reached, a trust mechanic designed through structured user research.
- WhatsApp and Telegram as the primary product surface, with LLM agents handling intake, matching, and follow-up so one person runs the day-to-day.
- Funnel and conversion analysis in SQL; pipeline instrumented to surface drop-off points and optimise routing logic.
- Directly relevant to Liquidity's agent-based origination and monitoring workflows: same pattern of automating a high-touch, relationship-dependent process without losing the human judgment layer.

§ HOW I WOULD WORK AT LIQUIDITY

First weeks, concrete.

WEEK 1

Listen, map, find the gaps.

Shadow the origination, diligence, and monitoring workflows end to end, interview analysts, data scientists, and engineers to map where the current AI outputs break down, then draft a ranked list of reliability and explainability gaps.

WEEK 2

Instrument, slice, validate.

Pick the single highest-leverage gap from week one, define an evaluation metric for it, and run a targeted prompt or RAG experiment with the data science team to confirm whether it is a data, model, or UX failure before writing a line of product spec.

WEEK 3+

Ship, close loops.

Own the first production improvement end to end, from revised prompt or pipeline change through to user acceptance and a monitored rollout, then set the cadence for ongoing AI output quality review across origination and risk scoring.

I COVER LETTER

Dear Liquidity Hiring Team,

Private credit decision-making is an unusually demanding environment for AI systems: the inputs are messy and heterogeneous, the stakes are high enough that hallucinations have real financial consequences, and the workflows, origination, diligence, monitoring, are deeply human and relationship-driven. That combination of technical complexity and domain sensitivity is exactly where I have spent the last several years as a product manager, and it is why this role caught my attention.

On the FCAS AI and data platform at Rohde & Schwarz I owned an internal AI and data platform serving multinational engineering teams working on a programme where accuracy and auditability were non-negotiable. I built the GitOps delivery pipeline, instrumented competing workflows to find out which one analysts actually used, and redirected the product strategy accordingly, which grew pipeline runs by 200%. At the Bundesagentur für Arbeit I productised

the ETL layer with Dagster and dbt, cutting data update time by 75%, and carried a full container migration alongside a UI relaunch without dropping delivery cadence. In both cases the core skill was the same: translate a complex, high-stakes decision process into a system that is fast enough to be useful and reliable enough to be trusted.

The responsibilities in this posting, defining AI-native experiences including agent-based workflows and decision support tools, owning data pipelines through to application layers, debugging across data, model, and UX, are a direct match for how I have operated. I have designed multi-agent orchestration systems using n8n and file-queue-based agent loops, and I have built and deployed LLM evaluation loops with human-in-the-loop checkpoints as a standard part of my production release process. On my own fintech project, Kvitt Payment Solutions, I built a mobile group-based P2P payments product from zero to acquisition inside the girocard ecosystem, which gave me direct experience with the financial data mapping and structured product representation work the posting describes.

The one honest gap is private credit as a specific domain. I have not previously owned credit risk scoring or deal structuring workflows. What I bring instead is a strong pattern-match: I have productised high-stakes decision AI in defence, public policy, and insurance, all regulated, all accuracy-critical, all requiring the same explainability and reliability standards that private credit demands. The asset class is new; the problem shape is familiar.

I am currently based in Berlin and actively targeting Abu Dhabi and the UAE, with a willingness to relocate on short notice. I would welcome a conversation about how my background maps to the specific AI systems Liquidity is building.

Best regards, *Philipp Hilbert*

II CURRICULUM VITAE

§ BEARING

Product manager with 13 years of ownership in regulated, data-intensive environments, defence AI, public-sector analytics, insurance, fintech, and healthcare. I translate complex business processes into production AI systems by owning the full stack: ETL and data pipelines, model evaluation and prompt engineering, and user-facing decision tools. Equally at home writing SQL and defining evaluation metrics as I am facilitating a stakeholder alignment session or restructuring a delivery team.

§ WHAT I CAN DO

AI Products and Tooling: LLM agents, agentic workflows, multi-agent orchestration, prompt and context engineering, RAG, evaluation loops, hallucination mitigation, human-in-the-loop design, AI governance, Explainable AI (XAI), EU AI Act

Analytics and Data: ETL (extract, transform, load), data pipelines, Dagster, dbt, SQL, PostgreSQL, funnel and churn analysis, KPI definition, Matomo, MicroStrategy

Product Management: continuous discovery, product roadmap, prioritisation, MVP definition, PoC, user research, prototyping, go-to-market

Product Ownership: requirements engineering, user stories, backlog refinement, sprint planning, story-point estimation, velocity, UAT, Jira, Confluence, Figma, Miro

Agile and Scaled Delivery: Scrum, Kanban, SAgile, ART, PI planning, CI/CD, GitOps with ArgoCD, release gating, defect-to-fix as a product metric

Technical Environments: Python, SQL, PostgreSQL, MongoDB, Kotlin (Spring Boot), REST APIs, Kubernetes, Docker, OpenShift, Jenkins, GitLab, Azure DevOps, Kafka

Governance and Business: GDPR, Deputy IT Security Officer, identity management, budget ownership, ROI reasoning, stakeholder and C-level reporting

§ PROFESSIONAL EXPERIENCE

ROHDE & SCHWARZ · PRODUCT MANAGER · DEFENCE AI AND DATA PLATFORM 2024-2025

Internal AI and data platform for the FCAS European programme; cross-functional team of 15+ engineers and data scientists. GitOps rollout with ArgoCD cut deployment time from days to 30 minutes; workflow instrumentation grew pipeline runs 200% by prioritising developer-code-first over UI-led design.

BUNDESAGENTUR FÜR ARBEIT · PRODUCT MANAGER · PUBLIC-SECTOR ANALYTICS 2023-2024

National labour-market analytics platform. Deputy IT Security Officer. Introduced Dagster and dbt, saving 75% of time per data update; migrated bare-metal to containers and carried a full ETL migration and UI relaunch in parallel.

CLINET PLATFORMS · PRODUCT OWNER · HEALTHCARE MOBILE PLATFORM 2023

Mobile patient app for iOS/Android. Cut roadmap to top 10% of features; user acceptance rose fivefold. Implemented CGM/KIS interface reducing digital anamnesis from 30 minutes to 5.

EMIL GROUP · PRODUCT MANAGER · INSURANCE B2B SAAS PLATFORM 2022

Restructured 18-person delivery org into three lean teams; attrition fell from near-total exodus to one departure. Shipped Claims Center from concept to production; cut new insurance product time-to-market from months to under a week.

SCHEIDT & BACHMANN · PRODUCT OWNER INTEGRATION · RETAIL PAYMENTS AND POS 2021-2022

Moved research results into production across a matrix organisation in a fully SAFe setup with agile release trains and PI planning.

BUNDESDRUCKEREI · PRODUCT OWNER (INTERIM) · SECURE GOVERNMENT DATA PLATFORM 2021

Data structures and identity management for the Robert Koch Institute programmes DIM and DESH.

AOK ITSCARE · PRODUCT OWNER (PROXY) · HEALTH INSURANCE B2B PROCUREMENT 2019-2020

Migrated legacy system to Shopware for three AOK insurers and ~20,000 employees; still in production. Backlog refinement, user stories and acceptance criteria in Jira.

ENERCON · PROJECT LEAD · WIND ENERGY SITE IT MODERNISATION 2019

Full site IT modernisation end to end.

WALDBURG-ZEIL KLINIKEN · PRODUCT OWNER · CLINICAL NETWORK INFRASTRUCTURE 2017-2018

Wifi infrastructure rollout across nine hospitals.

§ FOUNDER & EARLY-STAGE EXPERIENCE

CONSTRUCT8 · FOUNDER AND PRODUCT MANAGER · CONSTRUCTION STAFFING MARKETPLACE 2025-PRESENT

Zero-to-one B2B marketplace; LLM agents across the full funnel; anonymity-until-payment trust mechanic; funnel analysis in SQL.

QCRYPT AG · CO-FOUNDER AND PRODUCT OWNER · QUANTUM-SECURE ENCRYPTION 2016-2021

Three-layer hardware/software product installing in under two minutes; later Chair of the board.

SMART SOIL TECHNOLOGIES · CO-FOUNDER, CDO AND CFO · AGRI-
NANOTECHNOLOGY 2018-2021

Lab to investor rounds; owned financial planning and investor reporting.

KVITT PAYMENT SOLUTIONS · FOUNDER AND CFO · FINTECH P2P
PAYMENTS 2013-2018

Mobile group-based P2P payments product built from zero; acquired and operated by Sparkasse inside the girocard ecosystem.

§ LANGUAGES

German, native. English, C2.

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

TRUE NORTH · APPLICATION PACKAGE · TALENTMATE

→