

Philipp Hilbert. *AI Product Manager, investment intelligence at the frontier.*

Positioning, product manager with 13 years shipping regulated, data-intensive products who now builds AI-native systems end to end, from ETL pipelines to agent-based decision workflows.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear Liquidity Hiring Team,

Private credit decision-making is an unusually demanding environment for AI systems: the inputs are messy and heterogeneous, the stakes are high enough that hallucinations have real financial consequences, and the workflows, origination, diligence, monitoring, are deeply human and relationship-driven. That combination of technical complexity and domain sensitivity is exactly where I have spent the last several years as a product manager, and it is why this role caught my attention.

On the FCAS AI and data platform at Rohde & Schwarz I owned an internal AI and data platform serving multinational engineering teams working on a programme where accuracy and auditability were non-negotiable. I built the GitOps delivery pipeline, instrumented competing workflows to find out which one analysts actually used, and redirected the product strategy accordingly, which grew pipeline runs by 200%. At the Bundesagentur für Arbeit I productised the ETL layer with Dagster and dbt, cutting data update time by 75%, and carried a full container migration alongside a UI relaunch without dropping delivery cadence. In both cases the core skill was the same: translate a complex, high-stakes decision process into a system that is fast enough to be useful and reliable enough to be trusted.

The responsibilities in this posting, defining AI-native experiences including agent-based workflows and decision support tools, owning data pipelines through to application layers, debugging across data, model, and UX, are a direct match for how I have operated. I have designed multi-agent orchestration systems using n8n and file-queue-based agent loops, and I have built and deployed LLM evaluation loops with human-in-the-loop checkpoints as a standard part of my production release process. On my own fintech project, Kvitt Payment Solutions, I built a mobile group-based P2P payments product from zero to acquisition inside the girocard ecosystem, which gave me direct experience with the financial data mapping and structured product representation work the posting describes.

The one honest gap is private credit as a specific domain. I have not previously owned credit risk scoring or deal structuring workflows. What I bring instead is a strong pattern-match: I have productised high-stakes decision AI in defence, public policy, and insurance, all regulated, all accuracy-critical, all requiring the same explainability and reliability standards that private credit demands. The asset class is new; the problem shape is familiar.

I am currently based in Berlin and actively targeting Abu Dhabi and the UAE, with a willingness to relocate on short notice. I would welcome a conversation about how my background maps to the specific AI systems Liquidity is building.

Best regards, *Philipp Hilbert*
