

Philipp Hilbert. *Platform product owner for AI-native developer tooling.*

Positioning, Product Manager with a decade of shipped AI, DevOps, and platform work – ready to own the Software Factory roadmap from day one.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

§ WHY 1&1

The Software Factory is solving exactly the right problem: not another point tool, but a governed, multi-tenant AI platform that every development team can rely on – with cost transparency and human-in-the-loop control built in from the start. The explicit focus on SDLC automation through CI/CD-integrated AI agents, central MCP servers, and Jira/Confluence toolchain alignment maps closely to work I have done at Rohde & Schwarz and the Federal Employment Agency, where the challenge was always making complex infrastructure legible and trustworthy across teams. The governance angle – KPI dashboards per project, workflow, and run level – is a product problem as much as an engineering one, and that boundary is where I operate best. *That is exactly the environment where my strengths compound.*

§ AI-NATIVE STACK

Daily tools, not buzzwords.

- **Claude Code**, in-repo refactors and requirement drafting
- **GitHub Copilot**, inline completion and test generation
- **GPT**, spec drafting and edge-case generation
- **Cursor / VSCode**, LLM-assisted development
- **n8n**, multi-agent workflow orchestration
- **Lovable**, rapid front-end prototyping
- **Framer**, design-heavy prototype iteration
- **Gamma.app**, stakeholder presentation drafts

§ TECHNICAL ENVIRONMENT

Shipped with, not just listed.

Python, Kotlin (Spring Boot), JavaScript, PostgreSQL, MongoDB, REST APIs, Kubernetes, ArgoCD, Argo Workflows, Helm, OpenShift, Docker, Jenkins, Bitbucket, Git, GitOps, dbt, Dagster, DuckDB, Redpanda, Django, MicroStrategy, Jira, Confluence, CI/CD pipelines.

§ SIDE PROJECT

DEV AGENT FRAMEWORK [GITHUB.COM/HILBERTP](https://github.com/HilbertP) →

A personal agentic workflow framework for crypto market analysis and algorithmic decision-making.

- Orchestrates autonomous agents across data ingestion, signal generation, and execution decision layers.
 - Built on a file-queue-based agent loop pattern with modular task handoff between agents.
 - Integrates with market analytics tools and on-chain data sources for real-time execution context.
 - Direct proof of capability for the Software Factory's core concept: governed, observable, multi-step agent automation.
-

§ HOW I WOULD WORK AT 1&1

First weeks, concrete.

WEEK 1

Listen, map, find the gaps.

Shadow the AI/ML, DevOps, and platform teams; map the current state of the workflow library, agent configurations, and MCP server integrations against what the job description calls for – identify the first backlog items that are blocking teams from adopting the platform.

WEEK 2

Slice the backlog.

Run a structured prioritisation session with engineering and governance stakeholders to rank the first sprint's backlog against the dimensions the posting names explicitly: SDLC coverage, Human-in-the-Loop gate design, and KPI dashboard completeness.

Ship and unblock.

Deliver the first refined epics with acceptance criteria, coordinate rollout with the DevOps and CI/CD teams, and begin the enablement cycle so that development teams can onboard to the platform without hand-holding.

I COVER LETTER

Dear Hiring Team at 1&1,

The Software Factory role is the clearest product challenge I have seen in this market: take AI tooling that today every team builds independently and often badly, centralise it into a governed, multi-tenant platform, and make it genuinely useful across an organisation at scale. I have spent the past three years working on adjacent problems – and I would like to explain concretely why I believe I can own this roadmap from day one.

At Rohde & Schwarz I was responsible for a central AI and data management platform inside the European defence programme FCAS. The core product challenge was not the AI itself but everything around it: reproducible data transformation pipelines, audit trails, data provenance, SafeAI and ExplainableAI evaluation concepts – in short, the governance layer that makes AI outputs trustworthy for regulated, multinational stakeholders. I delivered autonomous feature sets end-to-end, from user stories and acceptance criteria through implementation support to GitOps-based rollout with ArgoCD. One infrastructure issue that the project lead had been requesting from engineers for over a year I unblocked and delivered independently in one week. That is the operational tempo and ownership depth I bring to a platform product role.

At the Federal Employment Agency I owned the Arbeitsmarktmonitor, an analytics platform with complex ETL pipelines, legacy system migrations, and a Dagster-based automation layer for indicator generation. I also served as Deputy IT Security Officer, which means I understand the compliance, data lineage, and governance requirements that the Software Factory must satisfy – including the EU data sovereignty constraints the job description names. The KPI and cost dashboard work described in the posting is familiar ground: I designed monitoring and transparency tooling for stakeholders who needed visibility into system behaviour without direct access to the underlying infrastructure.

My day-to-day product execution is AI-native by default. I use Claude Code for in-repo work and requirement generation, GitHub Copilot for inline development, and n8n for multi-agent workflow orchestration. I have built my own agentic framework for autonomous decision pipelines, which gives me a practitioner's understanding of where agent loops break, how human-in-the-loop gates need to be designed to be genuinely useful rather than just compliant, and what governance tooling actually needs to surface to be actionable. The Software Factory's

model – shared platform, central MCP servers, Jira and Bitbucket integration – is the architecture I would have proposed.

Earlier in my career I rebuilt delivery stability at Emil Group's InsurTech platform by splitting a bloated 17-person team, reintroducing sprint discipline and velocity tracking, and running weekly product workshops to rebuild client trust in a prototype-stage environment. Rollout, enablement, and change management – the last bullet in the job description – are things I have done under pressure, not just read about. I am comfortable at the boundary between strategic roadmap ownership and hands-on delivery support, and I operate well in organisations where both are required simultaneously.

I would welcome the opportunity to talk through how the Software Factory's current backlog maps to the platform maturity model I would bring to the role.

Best regards, *Philipp Hilbert*

II CURRICULUM VITAE

§ BEARING

Product Manager with nearly ten years of experience building AI systems, analytics platforms, and developer tooling in regulated and cross-functional environments. Delivered a central AI data management platform for a European defence programme, introduced GitOps-based automation that reduced deployment times from days to under an hour, and rebuilt delivery stability on a multi-team InsurTech platform under significant client pressure. AI-native by default: daily use of LLM-assisted development, multi-agent orchestration, and rapid prototyping. Equally comfortable writing an epic, unblocking an infrastructure issue, and presenting a governance model to senior stakeholders.

§ CORE STRENGTHS

- Multi-tenant AI platform product ownership and roadmap prioritisation
- SDLC automation: CI/CD, GitOps, agent-driven workflow design
- Governance, compliance, and data lineage in regulated environments
- Cross-functional alignment across AI/ML, DevOps, platform, and security teams
- Backlog management, sprint discipline, and delivery velocity recovery
- AI-native execution: LLM-assisted development, multi-agent orchestration

- Translating AI capability into measurable business and operational value
 - Human-in-the-loop process design and change management
 - Founder-level agency – low ego, high ownership, comfortable in ambiguity
-

§ PROFESSIONAL EXPERIENCE

ROHDE & SCHWARZ 2024, 2025

Product Manager, AI Data Transformation

- Delivered a central AI and data management platform for multimodal datasets within FCAS, covering data transformation, audit trails, and full data provenance
- Enabled ML training and evaluation with transparent SafeAI and ExplainableAI evaluation concepts
- Reduced deployment times from several days to 20–40 minutes by introducing GitOps with ArgoCD
- Delivered autonomous feature sets end-to-end: user stories, acceptance criteria, implementation support, testing, and GitOps-based rollout
- Unblocked an infrastructure-level OpenShift issue independently; delivered in one week a feature engineers had not shipped in over a year
- Led a cross-functional team including architect, data scientists, backend and frontend engineers, and QA

FEDERAL EMPLOYMENT AGENCY (BUNDESAGENTUR FÜR ARBEIT) 2023, 2024

Product Manager, Public Data Analytics

- Owned the Arbeitsmarktmonitor analytics platform for regional labour market and demographic data
- Introduced Dagster for automated ETL and indicator generation; standardised CI/CD with full test integration at build time
- Migrated legacy systems to Kubernetes and managed major system upgrades
- Delivered labour shortage analytics, flexible regional comparison tooling, and modernised admin interfaces
- Served as Deputy IT Security Officer with responsibility for compliance, governance, and data lineage
- Designed interactive dashboards and community features for a broad public-sector stakeholder base

EMIL GROUP 2022, 2022

Product Manager, InsurTech SaaS Platform

- Owned the product backlog for a B2B InsurTech platform reducing insurance time-to-market from over one year to under one week
- Delivered modules spanning product configuration, pricing, underwriting, policy issuance, document management, and workflow automation
- Designed and delivered a Claims Center module from concept to production rollout
- Re-established sprint discipline: split a bloated 17-person team into smaller Scrum units, introduced ownership concepts and deputy structures
- Reintroduced velocity tracking and structured refinement; reduced senior developer attrition from near-total to a single departure
- Rebuilt client trust in a prototype-stage environment through structured weekly product workshops

CLINET PLATFORMS 2021, 2022

Product Owner, Mobile Healthcare App

- Built a mobile-first pilot app for hospitals across iOS and Android, covering digital anamnesis, meal and therapy plans, transport, chat, and patient document storage
- Integrated CGM interface into hospital information system workflows
- Led a small team under tight deadlines, maintaining transparent stakeholder communication around legacy API constraints and external database dependencies

§ FOUNDER & EARLY-STAGE EXPERIENCE

QCRYPT AG 2016, 2018

Product Owner

- Owned the product architecture for a quantum-secure encryption product based on one-time-pad keys and true random number generator hardware
- USP: setup under two minutes, no technical knowledge required; served as deputy IT security officer

KVITT PAYMENT SOLUTIONS 2013, 2018

Founder & CFO

- Founded a group-based P2P payment product on a P2P backbone; owned PSP integration, UX design, and mobile app prototype
- Later acquired and operated by Sparkasse

Proxy Product Owner

- Migrated internal webshop to Shopware, including complex catalogue and API integration and configurable approval policies

§ LANGUAGES

German, native. English, C2.

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)