

Philipp Hilbert. *Platform product owner for AI-native developer tooling.*

Positioning, Product Manager with a decade of shipped AI, DevOps, and platform work – ready to own the Software Factory roadmap from day one.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear Hiring Team at 1&1,

The Software Factory role is the clearest product challenge I have seen in this market: take AI tooling that today every team builds independently and often badly, centralise it into a governed, multi-tenant platform, and make it genuinely useful across an organisation at scale. I have spent the past three years working on adjacent problems – and I would like to explain concretely why I believe I can own this roadmap from day one.

At Rohde & Schwarz I was responsible for a central AI and data management platform inside the European defence programme FCAS. The core product challenge was not the AI itself but everything around it: reproducible data transformation pipelines, audit trails, data provenance, SafeAI and ExplainableAI evaluation concepts – in short, the governance layer that makes AI outputs trustworthy for regulated, multinational stakeholders. I delivered autonomous feature sets end-to-end, from user stories and acceptance criteria through implementation support to GitOps-based rollout with ArgoCD. One infrastructure issue that the project lead had been requesting from engineers for over a year I unblocked and delivered independently in one week. That is the operational tempo and ownership depth I bring to a platform product role.

At the Federal Employment Agency I owned the Arbeitsmarktmonitor, an analytics platform with complex ETL pipelines, legacy system migrations, and a Dagster-based automation layer for indicator generation. I also served as Deputy IT Security Officer, which means I understand the compliance, data lineage, and governance requirements that the Software Factory must satisfy – including the EU data sovereignty constraints the job description names. The KPI and cost dashboard work described in the posting is familiar ground: I designed monitoring and transparency tooling for stakeholders who needed visibility into system behaviour without direct access to the underlying infrastructure.

My day-to-day product execution is AI-native by default. I use Claude Code for in-repo work and requirement generation, GitHub Copilot for inline development, and n8n for multi-agent workflow orchestration. I have built my own agentic framework for autonomous decision

pipelines, which gives me a practitioner's understanding of where agent loops break, how human-in-the-loop gates need to be designed to be genuinely useful rather than just compliant, and what governance tooling actually needs to surface to be actionable. The Software Factory's model – shared platform, central MCP servers, Jira and Bitbucket integration – is the architecture I would have proposed.

Earlier in my career I rebuilt delivery stability at Emil Group's InsurTech platform by splitting a bloated 17-person team, reintroducing sprint discipline and velocity tracking, and running weekly product workshops to rebuild client trust in a prototype-stage environment. Rollout, enablement, and change management – the last bullet in the job description – are things I have done under pressure, not just read about. I am comfortable at the boundary between strategic roadmap ownership and hands-on delivery support, and I operate well in organisations where both are required simultaneously.

I would welcome the opportunity to talk through how the Software Factory's current backlog maps to the platform maturity model I would bring to the role.

Best regards, *Philipp Hilbert*
