

# Philipp Hilbert. *AI-native Product Manager for Generative and Agentic systems.*

*Positioning*, ten years owning complex AI and data products end-to-end, now focused on enterprise AI that creates measurable business impact.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · [www.true-north.berlin](https://www.true-north.berlin) · [YouTube \(EN sample\)](#)

---

To the hiring team at AI Futures,

I am applying for the AI Product Owner role you are placing at a major German DAX 40 enterprise. My background is a close fit: nearly ten years as a hands-on Product Manager and Product Owner across AI data platforms, analytics systems, and regulated B2B SaaS, including direct ownership of Generative AI and data transformation products in a government-contract setting. I speak German natively, hold EU citizenship, and can relocate to Cologne.

The most directly relevant engagement is my recent role at Rohde & Schwarz, where I owned the product vision and delivery for a central AI and data management platform within the European defence programme FCAS. The platform handled multimodal datasets and fed curated data into ML training and evaluation pipelines. I introduced SafeAI and ExplainableAI evaluation concepts, built audit trails with full data provenance, and delivered autonomous feature sets end-to-end – from user stories through to GitOps rollout with ArgoCD. I reduced deployment times from several days to under forty minutes and unblocked an infrastructure-level OpenShift issue independently that the engineering team had been unable to resolve for over a year. That combination of product rigour and technical self-sufficiency is what I would bring to this role.

At the Bundesagentur für Arbeit I owned a public-facing analytics platform for regional labour market data. I introduced Dagster for automated ETL, managed a Kubernetes migration, served as Deputy IT Security Officer, and standardised CI/CD. This is the kind of large-organisation, compliance-adjacent environment the posting describes – and I have a working understanding of how to navigate governance requirements without letting them become a delivery bottleneck.

At Emil Group I inherited a broken delivery programme: a 17-person team with no sprint discipline, high senior developer attrition, and a client relationship close to collapse. I split the team into effective Scrum units, reintroduced velocity tracking and structured refinement, ran weekly product workshops to rebuild client trust, and reduced senior developer attrition from

near-total to a single departure. Stakeholder management in a high-pressure enterprise context is not theoretical for me – I have done it when the stakes were high.

I build and use AI-native tooling daily: I use Claude and GPT for requirements synthesis and edge-case generation, GitHub Copilot and Claude Code for in-repo acceleration, n8n for agentic workflow orchestration, and Lovable for rapid front-end prototyping in early discovery. I also maintain a personal backtesting and agent framework for algorithmic crypto strategy development, which has deepened my practical understanding of multi-agent orchestration architectures – the same patterns that underpin enterprise agentic AI initiatives.

I would welcome a conversation about the specific AI initiatives this enterprise is running and how my background maps to their current stage.

Best regards, *Philipp Hilbert*

---