

Philipp Hilbert. *Senior Product Manager, Agentic AI Assistants.*

Positioning, Product manager with 13 years shipping regulated B2B SaaS who builds agentic AI systems hands-on, evals, tool-calling, orchestration, and all.

BERLIN, GERMANY · HILBERT@TRUE-NORTH.BERLIN · TRUE-NORTH.BERLIN
LinkedIn · GitHub · www.true-north.berlin · YouTube (EN sample)

Dear Hiring Team,

Deel is doing something genuinely hard: turning the most compliance-dense corner of global business, cross-border payroll, worker classification, benefits, into a product experience that feels simple. The AI assistants role is the place where that ambition either becomes real or stays a demo, and that is the kind of inflection point I want to be at.

I have spent 13 years shipping product in regulated, complexity-heavy domains: a defence AI and data platform for the FCAS programme at Rohde & Schwarz, a national labour-market analytics platform at the Federal Employment Agency, clinical data integration at CLINET, an insurance SaaS at EMIL Group. In every case the challenge was the same: deeply structured domain knowledge, high stakes if something is wrong, and users who will not tolerate an answer that is merely plausible. That framing maps directly onto what a Deel AI assistant needs to be.

The hands-on piece matters to me, and I want to be direct about it. At Rohde & Schwarz I did not specify the AI platform; I prototyped agent workflows, instrumented two competing interaction patterns, and used the traces to kill the one that looked good but got zero real usage. At construct8, my construction staffing marketplace, I built the entire agentic funnel myself, LLM-driven intake, screening, and matching on WhatsApp and Telegram, with n8n orchestrating the workflow and SQL measuring where the funnel broke. I have written evals, designed tool-calling schemas, and debugged retrieval failures. I use Claude Code and Cursor daily. When you say "be hands-on with the technology", I take that to mean what I already do.

What I would bring to the Agentic AI Assistants role specifically: a habit of designing evaluation systems before writing a single line of agent code, a bias toward shipping small and measuring rather than planning large, and the cross-functional credibility, with engineers, data scientists, legal, and support, to keep a complex AI product moving without turning it into a process machine. I have restructured a failing 18-person delivery org at EMIL Group, introduced GitOps rollout at Rohde & Schwarz that cut deployment time from days to 30 minutes, and watched a beautiful UI die because I instrumented it. I make decisions from data, including decisions to kill my own ideas.

I am an EU citizen based in Berlin, available on four weeks' notice, and the fully remote structure is a natural fit. I would be glad to walk through a concrete eval design or agent architecture problem in a working session, that tends to be a faster signal than a further round of interviews.

Best regards, *Philipp Hilbert*
