

APPLICATION ·
SIXT

(SENIOR) TECHNICAL PRODUCT OWNER - AI PLATFORM (M/
F/D)

Philipp Hilbert. *Platform-minded Product Owner for AI infrastructure.*

Positioning, ten years owning technical backlogs in regulated, data-heavy environments – now applying that depth to MLOps and LLM platform delivery at scale.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

§ WHY SIXT

SIXT is building a central AI platform that serves over 80 data scientists and ML engineers – a scale where platform product ownership stops being abstract and starts requiring real prioritization under competing demands from many technical users. The job description explicitly names LLM Gateway, RAG, observability, and agentic frameworks, which maps precisely to the infrastructure layer I find most technically interesting. What draws me further is the dual-product scope: owning both the MLOps pipeline layer and the AI Service tooling simultaneously is the kind of complex, interdependent challenge that rewards systems thinking over checklist management. *That is exactly the environment where my strengths compound.*

§ HONEST FRAMING

My direct MLOps/LLMOps platform work is primarily from the product side of a defence-grade AI data management system, not from operating a commercial inference or model-serving fleet at the scale SIXT runs. The transferable mechanic is strong, though: I have repeatedly owned the prioritization layer between data scientists, ML engineers, and platform infrastructure, translated their competing needs into a coherent backlog, and shipped infrastructure-level changes end-to-end.

§ AI-NATIVE STACK

Daily tools, not buzzwords.

- **Claude**, spec generation and edge-case reasoning
- **GPT**, research synthesis and stakeholder document drafting
- **GitHub Copilot**, in-editor implementation support
- **Claude Code**, in-repo refactors and acceptance criteria generation
- **Cursor / VSCode**, day-to-day development environment
- **Lovable**, rapid front-end prototyping
- **n8n**, agentic workflow automation and multi-agent orchestration
- **Gamma.app**, fast stakeholder presentation drafts

§ TECHNICAL ENVIRONMENT

Shipped with, not just listed.

Kubernetes, ArgoCD, Argo Workflows, Helm, OpenShift, Docker, Jenkins, GitOps pipelines, Python, Kotlin (Spring Boot), PostgreSQL, MongoDB, Redpanda, dbt, Dagster, DuckDB, Django, JavaScript, REST APIs, CI/CD with full test integration at build time.

§ SIDE PROJECT

BACKTESTING ENGINE & AGENT FRAMEWORK [GITHUB.COM/HILBERTP](https://github.com/HilbertP) →

Personal infrastructure project building a crypto market backtesting engine alongside a file-queue-based autonomous dev-agent framework.

- Implements a backtesting engine for crypto markets covering perpetual futures, spot mechanics, and execution quality analysis.
- Architecture uses a file-queue-based agent loop for autonomous task orchestration, directly analogous to agentic pipeline design on LLMOps platforms.
- Integrates market microstructure analytics including slippage, price impact, liquidation heatmaps, and routing logic.
- Demonstrates the systems-building and infrastructure-first thinking directly applicable to owning an MLOps/LLMOps platform backlog.

§ HOW I WOULD WORK AT SIXT

First weeks, concrete.

WEEK 1

Listen, map, find the gaps.

Shadow the data scientists, ML engineers, and platform engineers who use both products daily; map the current pain points in the model lifecycle and LLM Gateway, and identify where the backlog diverges from actual user need.

WEEK 2

Slice, validate, prioritize.

Run a structured refinement session with the team to size and sequence the top five backlog items against adoption impact, infrastructure risk, and dependency chains across both platform products.

WEEK 3+

Ship, align, close.

Drive the first sprint to a clean review with clear acceptance criteria met, establish a lightweight stakeholder cadence for the international audience, and own the tradeoff communication when priorities conflict.

I COVER LETTER

Dear SIXT Hiring Team,

The role of Technical Product Owner for the MLOps/LLMOps Platform and AI Service is one I can step into with immediate relevance. You are running a central platform that 80-plus data scientists and ML engineers depend on – which means the cost of a poorly prioritized backlog, an unresolved dependency, or a stalled technical discussion is multiplied across every downstream team. I have operated in exactly that accountability structure, most recently at Rohde & Schwarz where I owned a central AI and data management platform for the FCAS defence programme.

At Rohde & Schwarz I held the full product ownership lifecycle for a multimodal data platform: data ingestion, transformation pipelines, audit trails with full provenance, and SafeAI and ExplainableAI evaluation concepts. I delivered autonomous feature sets end-to-end – from user stories through implementation support to GitOps-based rollout via ArgoCD – and reduced

deployment times from several days to 20–40 minutes by introducing GitOps into a team that had not previously operated that way. I also unblocked an infrastructure-level OpenShift issue independently that the project lead had been requesting from engineers for over a year; I delivered it in one week. When you describe needing someone who can moderate technical discussions and drive the team to clear decisions, that is a direct description of how I operated there.

At Emil Group I inherited a broken delivery environment: a 17-person team with no sprint discipline, near-total senior developer attrition risk, and a client who had lost confidence in the product. I re-established Scrum structure, split the team into functional units with clear ownership, reintroduced velocity tracking, and regained client trust through structured weekly product workshops. This experience sharpened my ability to hold both the technical prioritization and the stakeholder communication simultaneously without letting either degrade.

Your job description calls out AWS SageMaker and Bedrock, LLM gateways, RAG, and agentic frameworks including LangGraph and LangChain as desirable. My LLM gateway and RAG work is primarily from the user and product side rather than from operating those services as a platform engineer, and I name that honestly. What I bring is a strong systems-thinking foundation, hands-on familiarity with the infrastructure layer those services sit on (Kubernetes, Argo Workflows, CI/CD), and a track record of translating data scientists' and ML engineers' actual needs into a backlog that ships. I learn infrastructure domains quickly because I engage with them technically rather than abstractly.

I am based in Berlin, open to relocating to Munich, and available with four weeks notice. I would welcome a conversation about where the platform stands and what the next six months of roadmap look like.

Best regards, *Philipp Hilbert*

II CURRICULUM VITAE

§ BEARING

Product Manager and Product Owner with nearly ten years owning technically complex backlogs across AI data platforms, analytics infrastructure, and regulated SaaS environments. Experienced in translating data scientists' and ML engineers' operational needs into sequenced, shipped platform capabilities. Comfortable holding both technical prioritization and international stakeholder alignment in parallel. Delivered end-to-end across Kubernetes-native

environments with GitOps, Argo Workflows, and full CI/CD discipline. Low ego, high agency, direct communicator.

§ CORE STRENGTHS

- MLOps/LLMOps platform product ownership and backlog prioritization
 - AI and data pipeline delivery: provenance, audit, transformation, evaluation
 - Cross-functional alignment across data science, engineering, and QA
 - Sprint discipline recovery and delivery cadence ownership
 - Technical stakeholder moderation and decision driving
 - GitOps and Kubernetes-native deployment ownership
 - Autonomous feature delivery end-to-end: stories to rollout
 - Risk-first discovery under ambiguity and incomplete requirements
 - Founder-level pragmatism and dependency management
-

§ PROFESSIONAL EXPERIENCE

ROHDE & SCHWARZ 2024, 2025

Product Manager, AI Data Transformation – FCAS (Future Combat Air System)

- Owned the product backlog for a central AI and data management platform serving multimodal datasets (image, audio, tabular) for ML training and evaluation
- Delivered autonomous feature sets end-to-end: user stories, acceptance criteria, implementation support, testing, and GitOps-based rollout via ArgoCD
- Reduced deployment times from several days to 20–40 minutes by introducing GitOps across the team
- Unblocked infrastructure-level OpenShift issues independently; delivered in one week a feature the project lead had been requesting from engineers for over a year
- Enabled transparent ML evaluation concepts including SafeAI and ExplainableAI with full data provenance and audit trails
- Led cross-functional team of architect, data scientists, backend and frontend engineers, and QA across multinational stakeholders

Tech: Kotlin (Spring Boot), Python, PostgreSQL, MongoDB, Redpanda, Kubernetes, Argo Workflows, ArgoCD, OpenShift

FEDERAL EMPLOYMENT AGENCY (BUNDESAGENTUR FÜR ARBEIT) 2023, 2024

Product Manager, Public Data Analytics – Arbeitsmarktmonitor

- Introduced Dagster for automated ETL and indicator generation pipelines across a regional labour market analytics platform
- Delivered new analytics features including labour shortage indicators, flexible regional comparison, and modernised admin tooling
- Migrated legacy systems to Kubernetes and managed major system upgrades with standardised CI/CD and full test integration
- Served as Deputy IT Security Officer (IT-SV) with responsibility for compliance, governance, and data lineage
- Designed interactive dashboards and community features serving public and institutional stakeholders

Tech: Python, Django, JavaScript, PostgreSQL, dbt, Dagster, DuckDB, MicroStrategy, Kubernetes, Jenkins, Bitbucket, Git

EMIL GROUP 2022, 2022

Product Manager, InsurTech SaaS Platform

- Re-established sprint discipline and delivery stability in a stalled environment; split a 17-person team into focused Scrum units with clear ownership and deputy structures
- Delivered a Claims Center module from concept to production rollout on a B2B insurance configuration platform
- Reintroduced velocity tracking and structured refinement after practices had been lost under delivery pressure
- Regained client trust through structured weekly product workshops in a prototype-stage environment
- Reduced senior developer attrition from near-total to a single departure through improved team clarity and process

Tech: Low-code platform, Java, HTML/CSS, Kubernetes, REST APIs

CLINET PLATFORMS 2021, 2022

Product Owner, Mobile Healthcare App

- Delivered core hospital app flows covering digital anamnesis, meal and therapy plans, transport, chat, and patient document storage on iOS and Android

- Integrated CGM interface into hospital information system workflows with legacy API constraints
- Built stakeholder trust through transparent communication of technical complexity across external database and legacy API dependencies

Tech: AngularJS, Ionic, Python, Kubernetes, CGM API integration

§ FOUNDER & EARLY-STAGE EXPERIENCE

QCRYPT AG 2016, 2018

Co-Founder and Product Owner

- Owned product architecture for a quantum-secure encryption product based on one-time-pad keys and true random number generator hardware
- Designed for setup under two minutes with no technical knowledge required from end users
- Served as deputy IT security officer; responsible for attack vector mitigation and security product positioning

KVITT PAYMENT SOLUTIONS 2013, 2018

Founder and CFO

- Founded and built a group-based P2P payments product on a P2P backbone from concept through mobile app prototype
- Owned PSP integration, full UX/UI design, and payment edge-case handling
- Product later acquired and operated by Sparkasse

SMART SOIL TECHNOLOGIES 2016, 2021

Co-Founder, CDO, CFO, Head of HR (successive roles)

- Led digital tooling establishment, financial and investment planning, and ERP introduction across founding and scaling phases

§ LANGUAGES

German, native. English, C2.

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

