

Philipp Hilbert. *Platform-minded Product Owner for AI infrastructure.*

Positioning, ten years owning technical backlogs in regulated, data-heavy environments – now applying that depth to MLOps and LLM platform delivery at scale.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear SIXT Hiring Team,

The role of Technical Product Owner for the MLOps/LLMOps Platform and AI Service is one I can step into with immediate relevance. You are running a central platform that 80-plus data scientists and ML engineers depend on – which means the cost of a poorly prioritized backlog, an unresolved dependency, or a stalled technical discussion is multiplied across every downstream team. I have operated in exactly that accountability structure, most recently at Rohde & Schwarz where I owned a central AI and data management platform for the FCAS defence programme.

At Rohde & Schwarz I held the full product ownership lifecycle for a multimodal data platform: data ingestion, transformation pipelines, audit trails with full provenance, and SafeAI and ExplainableAI evaluation concepts. I delivered autonomous feature sets end-to-end – from user stories through implementation support to GitOps-based rollout via ArgoCD – and reduced deployment times from several days to 20–40 minutes by introducing GitOps into a team that had not previously operated that way. I also unblocked an infrastructure-level OpenShift issue independently that the project lead had been requesting from engineers for over a year; I delivered it in one week. When you describe needing someone who can moderate technical discussions and drive the team to clear decisions, that is a direct description of how I operated there.

At Emil Group I inherited a broken delivery environment: a 17-person team with no sprint discipline, near-total senior developer attrition risk, and a client who had lost confidence in the product. I re-established Scrum structure, split the team into functional units with clear ownership, reintroduced velocity tracking, and regained client trust through structured weekly product workshops. This experience sharpened my ability to hold both the technical prioritization and the stakeholder communication simultaneously without letting either degrade.

Your job description calls out AWS SageMaker and Bedrock, LLM gateways, RAG, and agentic frameworks including LangGraph and LangChain as desirable. My LLM gateway and RAG work is primarily from the user and product side rather than from operating those services as a platform engineer, and I name that honestly. What I bring is a strong systems-thinking foundation, hands-on familiarity with the infrastructure layer those services sit on (Kubernetes, Argo Workflows, CI/CD), and a track record of translating data scientists' and ML engineers' actual needs into a backlog that ships. I learn infrastructure domains quickly because I engage with them technically rather than abstractly.

I am based in Berlin, open to relocating to Munich, and available with four weeks notice. I would welcome a conversation about where the platform stands and what the next six months of roadmap look like.

Best regards, *Philipp Hilbert*
