

Philipp Hilbert. *AI Product Manager with 13 years shipping complex platforms.*

Positioning, I own AI product strategy end-to-end, from RAG and agentic workflows through evaluation pipelines to production governance.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear Open Innovation AI team,

The problem Open Innovation AI is solving, making AI workload management hardware-agnostic and operationally simple across diverse GPU and accelerator infrastructure, is one of the more underserved product challenges in the current AI landscape. Most AI PM roles sit on top of the stack: prompt, evaluate, ship a feature. OICM is the layer underneath, and owning the product strategy for an orchestration platform that enterprises depend on for cost, latency, and reliability is a different kind of responsibility. That is the work I want to do.

My 13 years of product ownership cover two things that are directly relevant here. First, I have shipped AI and data platforms in production, not as demos. At Rohde & Schwarz I owned the internal AI and data platform for the FCAS European programme, managing a cross-functional team of 15+ engineers and data scientists. Redirecting the product strategy based on actual workflow instrumentation, finding that the UI-led approach recorded zero clicks, drove a 200% increase in pipeline runs. At the Federal Employment Agency I owned an ETL-heavy national labour-market analytics platform, introduced Dagster and dbt to replace fragile hand-held imports, and saved roughly 75% of the time per data update. Both roles required translating ambiguous infrastructure decisions into clear product choices with measurable outcomes.

Second, I work directly with the AI-native stack the role describes. I design agentic workflows, write evaluation loops, think through RAG architectures and model selection tradeoffs, and use Claude Code, GitHub Copilot and Cursor as daily tools rather than occasional aids. My background in Kubernetes, ArgoCD, and GitOps delivery means I can sit with your engineering team at the architecture level, not just relay requirements from a distance.

The skills and experience required in the posting, LLMs and foundation models, RAG pipelines, agents and orchestration frameworks, cost optimisation across tokens and GPU infra, experiment design, evaluation and guardrails, match my current practice closely. I am also familiar with AI governance requirements including the EU AI Act, which matters for enterprise customers who need to justify workload decisions to their compliance teams.

I would be glad to talk through any of this in more detail. I am available at four weeks' notice and open to a remote arrangement.

Best regards, *Philipp Hilbert*
