

Philipp Hilbert. *Senior Product Owner for mission-critical digital platforms.*

Positioning, 13 years owning multi-disciplinary technical product backlogs in regulated environments, from defence AI to public-sector data platforms.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear DigyCorp Hiring Team,

The Senior Product Owner role for the Digital Twin Platform stood out immediately because it is structurally the hardest kind of product job: you need someone who can hold a rigorous, evidence-based conversation with senior research scientists about reef ecology, then turn around and translate that conversation into a backlog the engineering teams can plan a sprint against, without losing precision in either direction. That is the operating mode I have been in for the last several years.

At Rohde & Schwarz I was the product authority for an internal AI and data Platform supporting the FCAS European defence programme. The stakeholder surface spanned defence programme leads, data scientists and a multinational engineering team of 15+ people working across cloud, DevOps and data engineering. I ran dual-track continuous discovery to keep requirements validated ahead of sprint commitment, owned the full backlog with clear epics, user stories and acceptance criteria, and prioritised using cost-of-delay reasoning to sequence work that actually mattered to the programme. The result was a 200% increase in pipeline runs after redirecting strategy based on instrumented user workflow analysis, and a GitOps rollout path that compressed deployment from days to 30 minutes, making the Platform hotfix-capable for the first time.

At the Bundesagentur für Arbeit I owned the product backlog for a national labour-market analytics Platform serving policymakers and researchers. That environment required the same combination of scientific-stakeholder management and deep technical delivery: I carried a full ETL migration, a bare-metal-to-containers modernisation and a UI relaunch simultaneously, while introducing Dagster and dbt to eliminate the hand-holding that had been consuming 75% of the time on each quarterly data update. The Definition of Done on every data pipeline story included automated test coverage in the Jenkins CI/CD pipeline, because a data Platform that scientists rely on cannot ship untested.

The DTP's scope, 2 petabytes of scientific, environmental, spatial and operational data, machine learning across the environment, and operational tooling used daily by nursery and field teams, maps directly to what I have shipped. I am comfortable owning the backlog across cloud engineering, DevOps, backend and frontend development and data engineering, comfortable presenting trade-off recommendations to a Change Control Board, and comfortable running story-mapping sessions with domain experts who are not product people. I use Jira, Azure DevOps and Confluence daily, and I use AI-assisted tooling, Atlassian RoVo, Claude Code, GitHub Copilot, to keep backlog documentation, acceptance criteria and sprint reporting current without the overhead that typically falls behind in fast-moving programmes.

The onsite delivery model and the governance structure, two-week sprints, rolling quarterly plan, formal CCB, are the conditions I work best in. Delivery discipline is not a constraint I work around; it is the mechanism that makes ambitious long-term programmes like this one trustworthy to their scientific stakeholders and their funders. I would be glad to discuss the role and the programme in more detail.

Best regards, *Philipp Hilbert*
