

# Philipp Hilbert. *Enterprise AI product, shipped under sovereign constraint.*

*Positioning*, ten years building regulated-industry platforms, now applying that to agentic AI that has to work inside the security perimeter, not around it.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · [www.true-north.berlin](https://www.true-north.berlin) · [YouTube \(EN sample\)](#)

---

## § WHY AI71

ai71 is solving the version of the problem that everyone else defers: AI that runs inside the customer's stack, inside their language, inside their legal perimeter, with no dependency on your cloud. That constraint is not a limitation of the product, it is the product. The combination of agentic workflow design, eval ownership, and forward-deployed learning is exactly the structure I have been building toward: short feedback loops, real operational stakes, and no deck to hide behind. The fact that Ask is infrastructure-agnostic and model-agnostic is the right architectural call, and I would want to own part of making that promise hold up in production. *That is the kind of hard problem that sharpens rather than blunts over time.*

---

## § HONEST FRAMING

I have not shipped inside a sovereign cloud or air-gapped government network directly, my regulated-environment experience sits in defence, national health infrastructure, and public-sector analytics in Germany, all of which carry similar procurement weight and privacy requirements but not the same physical isolation. The transferable mechanic is the same: working through heavyweight compliance cycles, designing for environments where you cannot iterate freely after go-live, and treating auditability as a first-class product requirement rather than an afterthought.

---

## § AI-NATIVE STACK

*Daily tools, not buzzwords.*

- **Claude Code**, in-repo refactors and feature delivery
- **GitHub Copilot**, day-to-day coding assistance
- **Lovable**, full-stack prototype in a day
- **n8n**, multi-agent workflow orchestration
- **GPT**, spec generation and edge-case enumeration
- **Framer**, design-heavy validation prototypes
- **Cursor/VSCode**, primary development environment
- **Gamma.app**, rapid structured communication

---

## § TECHNICAL ENVIRONMENT

*Shipped with, not just listed.*

Python, Kotlin, Kafka, Dagster, dbt, PostgreSQL, MongoDB, REST APIs, Kubernetes, OpenShift, ArgoCD, Helm, Jenkins, GitOps, Docker, Argo Workflows, WhatsApp and Telegram messaging APIs, Shopware, CGM/KIS clinical interfaces, LLM agents, n8n, Lovable, Claude Code, GitHub Copilot.

---

## § SIDE PROJECT

DEV AGENT FRAMEWORK [GITHUB.COM/HILBERTP](https://github.com/hilbertp) →

*A personal agentic infrastructure for crypto market analysis and automated execution.*

- Builds and runs file-queue-based agent loops for backtesting crypto trading strategies across centralized and decentralized venues.
- Uses a modular agent orchestration pattern: discrete task agents coordinated by a lightweight queue broker rather than a monolithic pipeline.
- Integrates live data from Coinglass and Hyblock Capital for liquidation heatmap analysis and market microstructure signals.
- Demonstrates the same self-built eval discipline the role asks for: the backtest engine is the ground-truth judge before any strategy touches real capital.

---

## § HOW I WOULD WORK AT AI71

*First weeks, concrete.*

### WEEK 1

#### **Listen, map, find the gaps.**

Shadow forward-deployed colleagues inside at least one live customer environment, map the actual workflow the agent is running against the PRD that specified it, and identify where the gap between demo behaviour and production behaviour is widest.

### WEEK 2

#### **Build the eval.**

Take the single most disputed quality question from week 1, define what good looks like in writing, and ship a first test set and judge so the squad has a measurable bar before the next sprint closes.

### WEEK 3+

#### **Ship, close loops.**

Own one PRD end-to-end through to deployment, wire the eval into the release gate, and bring one learning from the forward-deployed channel back into the roadmap with a concrete proposal.

---

## **I** COVER LETTER

Dear ai71 team,

I am applying for the Product Manager role because the problem you are solving is the one I find most interesting in enterprise AI right now: making agentic systems trustworthy enough to run a ministry's legal review or a bank's procurement workflow, inside the customer's own perimeter, without any of the escape hatches that make demos look better than they are.

My most relevant experience for this role is at Rohde & Schwarz, where I owned the defence AI platform for the FCAS programme. The environment there was not air-gapped, but it carried the same procurement weight, classification sensitivity, and expectation of auditability that you describe. I moved that platform from a UI-heavy approach to a developer-code-friendly one, cut rollout time from days to thirty minutes, and made the system fully hotfix-capable, all without the freedom to iterate after go-live that a consumer product would allow. Before that, at Bundesdruckerei, I owned the national health data platform for the RKI, including the

vaccination statistics and viral-mutation tracking systems that ran under Germany's highest privacy requirements. Both roles required treating compliance and auditability as first-class product requirements, not layers bolted on at the end.

The eval ownership requirement in your job description is the piece I most want to address directly, because it is the one most PM roles avoid. At construct8, I own the quality bar for LLM-agent outputs end-to-end, defining what "correct" means, building test sets, and making the agent's behaviour measurable before the customer sees it. I built my own agent framework as a side project, including a backtesting engine that functions as a ground-truth judge before any strategy runs on real capital. The pattern is the same: you cannot ship a probabilistic system into a regulated environment and rely on "it looked fine in the demo." I share that view, and I know how to build the scaffolding that makes quality visible.

On the builder requirement: I prototype in Lovable, run multi-agent workflows in n8n, and ship feature-complete code via Claude Code and GitHub Copilot. A disputed flow becomes a working artefact the same day. I bring that to product debates rather than a slide deck, and I expect that to be useful in a squad that ships and learns faster than the roadmap predicts.

The forward-deployed track is where I would want to start. The shortest distance between a product decision and a real operation is exactly the feedback loop that makes the rest of the roadmap honest.

Best regards, *Philipp Hilbert*

---

## *II* CURRICULUM VITAE

---

### § BEARING

Ten years of product ownership in regulated environments, defence, national health infrastructure, public-sector analytics, InsurTech, where the cost of a wrong call is measured in procurement cycles and operational downtime rather than sprint velocity. I work AI-first: a disputed flow becomes a clickable prototype in a day, and I own the production code without writing it. My orientation is outcome over output, and I define success before we build so that whether it worked is never a matter of opinion.

---

### § WHAT I CAN DO

**Data and AI platforms.** Multimodal AI training data pipelines, ETL orchestration, LLM agents, eval design, agentic workflow orchestration. Kafka, Python, Kotlin, Dagster, dbt.

**Product craft.** Continuous discovery, prototype-led validation, outcome negotiation, backlog and release ownership, SAFe (PI planning, release trains), metric definition before build.

**AI-first building.** Claude Code, GitHub Copilot, Lovable, n8n, Cursor/VSCode, GPT, Framer, Gamma.app.

**Cloud, delivery and operations.** Bare-metal to container migration, GitOps, rollout and hotfix pipelines. Kubernetes, OpenShift, ArgoCD, Helm, Jenkins, Docker, Azure DevOps.

**Regulated environments.** Defence, public sector, health data, insurance: identity management, audit and governance cycles, highest privacy requirements.

**Integration and platforms.** Clinical interfaces (CGM/KIS), payment and forecourt systems, Shopware, messaging APIs (WhatsApp, Telegram).

---

## § PROFESSIONAL EXPERIENCE

### CONSTRUCT8 · PRODUCT MANAGER · CONSTRUCTION MARKETPLACE, B2B SAAS 2025 TO NOW

Zero-to-one agentic platform; LLM-agent quality owned end-to-end, eval sets built before customer exposure. WhatsApp/Telegram APIs, LLM agents.

### ROHDE & SCHWARZ · PRODUCT MANAGER · DEFENCE AI PLATFORM (FCAS), B2B SAAS 2024 TO 2025

Multimodal AI data homogenised into one training standard; rollout cut from days to 30 min, fully hotfix-capable; platform reoriented from UI-heavy to dev-code-friendly. Kafka, Kotlin, Python.

### BUNDESAGENTUR FÜR ARBEIT · PRODUCT MANAGER · PUBLIC ANALYTICS PLATFORM 2023 TO 2024

Public labour-market data made searchable, filterable, and social-media-ready; bare-metal-to-container migration and big ETL migration shipped. Dagster, dbt.

### EMIL GROUP · PRODUCT MANAGER · INSURTECH LOW-CODE SAAS 2022

Claims Center launched; 17-person monolithic team restructured into three nimble squads with leads; insurers launching products in days instead of months.

### CLINET · PRODUCT OWNER · HEALTH-TECH MOBILE APP 2021 TO 2022

Native iOS/Android app; CGM/KIS clinical interface integrated directly.

BUNDESDRUCKEREI · PRODUCT OWNER · NATIONAL HEALTH DATA  
PLATFORM (RKI) 2021

RKI projects DIM and DESH: vaccination statistics and viral-mutation tracking under  
Germany's highest privacy and identity-management requirements.

SCHEIDT & BACHMANN · PRODUCT OWNER · FUEL-RETAIL INTEGRATION  
2020 TO 2021

POS, payment, and forecourt systems integrated at scale. SAFe: PI planning, release trains.

AOK ITSCARE · PROXY PRODUCT OWNER · PROCUREMENT PLATFORM 2019  
TO 2020

Full legacy-to-Shopware migration for ~20,000 users delivered on time.

ENERCON · PROJECT LEAD · IT MODERNISATION, WIND ENERGY 2019

Full site modernisation: end devices, servers, network, software.

---

§ FOUNDER & EARLY-STAGE EXPERIENCE

Kvitt (mobile payments, acquired by Sparkasse) · Qcrypt (encryption, quantum-resistant OTP/  
true-RNG; deputy IT security officer) · Smart Soil Technologies (nanotech)

---

§ LANGUAGES

German, native. English, C2.

---

[LinkedIn](#) · [GitHub](#) · [www.true-north.berlin](http://www.true-north.berlin) · [YouTube \(EN sample\)](#)