

Philipp Hilbert. *AI-native Product Manager for enterprise workflow transformation.*

Positioning, ten years translating complex operational workflows into production software, now building AI-native products where the requirement is to reinvent the work itself.

BERLIN, GERMANY · HILBERT@TRUE-NORTH.BERLIN · TRUE-NORTH.BERLIN

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

§ WHY CHAPTER

CHAPTER sits inside a major publishing group but operates as a startup with a specific mandate: prove that generative AI can reshape how large organisations actually work, not just automate surface tasks. The job description calls out "unearthing new possibilities across the entire value chain" – that is exactly the kind of systems-level problem I have been solving, from defence AI data pipelines at Rohde & Schwarz to insurance workflow automation at Emil Group. A product culture that is remote-first, family-first, and reporting directly to the Director of Product removes the usual political drag and keeps the focus on delivery. *That is the structural setup where a high-agency PM compounds fastest.*

§ AI-NATIVE STACK

Daily tools, not buzzwords.

- **Claude**, spec drafting, edge-case generation, requirements review
- **GPT**, research synthesis and stakeholder communication drafts
- **Cursor / VSCode**, in-repo review and implementation support
- **Lovable**, rapid front-end prototyping before engineering commits
- **Framer**, design-heavy wireframing and click-through prototypes
- **Gamma.app**, internal product narratives and stakeholder decks

- **n8n**, multi-agent workflow orchestration and automation design
 - **GitHub Copilot**, developer pairing and acceptance-criteria validation
-

§ TECHNICAL ENVIRONMENT

Shipped with, not just listed.

Production work has touched Kotlin with Spring Boot, Python, Django, JavaScript, AngularJS, PostgreSQL, MongoDB, REST APIs, Redpanda, Kubernetes, ArgoCD, Argo Workflows, OpenShift, Helm, Dagster, dbt, DuckDB, MicroStrategy, Jenkins, Bitbucket, Docker, and low-code insurance platforms. Front-end prototyping via Lovable and Framer sits alongside that backend fluency.

§ SIDE PROJECT

BACKTEST ENGINE & DEV-AGENT FRAMEWORK [GITHUB.COM/HILBERTP](https://github.com/HilbertP) →

Personal crypto-market backtesting engine and agentic developer framework, built entirely outside working hours.

- Implements a file-queue-based agent loop for autonomous task execution across research, data ingestion, and signal generation.
 - Architecture separates the orchestration layer (n8n, custom Python agents) from the analytical core (DuckDB, time-series stores) for clean testability.
 - Active development with iterative commit cadence; trading strategy coverage spans perpetual futures, spot DEX execution, and cross-chain routing scenarios.
 - Demonstrates the end-to-end ownership CHAPTR needs: discovery, architecture decision, implementation, and iteration — all driven by one person with no external dependency.
-

§ HOW I WOULD WORK AT CHAPTR

First weeks, concrete.

WEEK 1

Listen, map, find the gaps.

Sit in on every live customer touchpoint and internal workflow demo I can access, map the current user journey end-to-end in writing, and identify the three places where the AI product either adds no value or actively creates friction.

WEEK 2

Slice and validate.

Take the highest-priority gap, write sharp problem statements and hypothesis cards, run two to three short user interviews or prototype tests to confirm or kill each assumption before engineering time is spent.

WEEK 3+

Ship and close loops.

Drive the first releasable increment from backlog-ready stories through engineering handoff to user feedback, then document the learning and fold it into the next roadmap slice.

I COVER LETTER

Dear CHAPTR team,

I have spent the past decade doing one thing in different settings: taking a messy, partially understood operational workflow and building the software product that makes it work at scale. At Rohde & Schwarz I did it inside a European defence programme, building a central AI data management platform for multimodal datasets from scratch. At Emil Group I did it inside B2B insurance, delivering a Claims Center module while simultaneously re-establishing sprint discipline for a 17-person team that had lost its delivery rhythm. At the Federal Employment Agency I did it for public labour market analytics, introducing Dagster-based ETL pipelines and modernising the platform to Kubernetes. The context changes; the core work does not.

CHAPTR's mandate is precisely that work, applied to the organisations that need it most: large enterprises trying to use generative AI to reshape daily workflows rather than bolt AI onto legacy processes. The job description asks for someone who can lead Product Discovery through Design Thinking, translate user needs into validated solutions, and drive iterations based on real feedback. That is my operating model. I run structured discovery workshops, write testable hypotheses, prototype fast in Lovable or Framer, and bring user interview evidence back to the team before engineering commits a single sprint. I have done this in regulated environments with demanding external stakeholders, which means I also know how to hold a product narrative steady when pressure comes from multiple directions at once.

I am genuinely AI-native, not as a talking point but as a daily working practice. I use Claude and GPT throughout the requirements and spec process, Cursor for in-repo review, and I have built a personal multi-agent orchestration framework in my own time to understand where these tools break under load. That background means I can have a credible technical conversation with the AI engineers on your team, push back on scope when the model is being asked to do something it cannot reliably do, and also protect the user experience when AI capability outpaces usability.

The specific pull of CHAPTR for me is the combination of startup execution pace inside a group that has real enterprise distribution. Springer Nature and Macmillan give the products a path to large-organisation customers; the startup structure means a PM can still make decisions at speed. Reporting directly to the Director of Product removes abstraction layers and puts accountability exactly where it belongs. I am a direct, low-ego operator who produces more in flat structures than in hierarchical ones, and the role description makes clear that is the culture you are building.

I would be glad to walk through any of the product decisions I have made in the past three years in a call. I can also show working prototypes and the agent framework I have built independently, which gives a clearer picture of how I actually work than a CV can.

Best regards, *Philipp Hilbert*

II CURRICULUM VITAE

§ BEARING

Hands-on Product Manager with nearly ten years of experience translating complex operational workflows into scalable, production-ready software. Background covers AI data platforms, B2B SaaS, regulated public-sector analytics, InsurTech, and founder-level company building. Fluent in AI-native working methods including LLM-assisted requirements engineering, rapid front-end prototyping, and autonomous feature delivery. Comfortable in ambiguity, precise in delivery, and effective as a cross-functional alignment point for engineering, design, data science, and business stakeholders. Ready to bring that full stack of skills to enterprise generative AI product development at CHAPTR.

§ CORE STRENGTHS

- Product discovery and Design Thinking facilitation with real users
- AI-native prototyping and requirements engineering
- Translating operational workflows into dev-ready product specs
- Cross-functional alignment across engineering, UX, data, and business
- Backlog management, sprint discipline, and velocity ownership
- KPI-oriented prioritisation and measurable iteration
- Workshop moderation, technical and non-technical audiences

- End-to-end feature delivery with GitOps and CI/CD fluency
 - Founder-level pragmatism, high agency, low ego
-

§ PROFESSIONAL EXPERIENCE

ROHDE & SCHWARZ 2024, 2025

Product Manager, AI Data Transformation

- Built a central AI and data management platform for multimodal datasets (image, audio, tabular) within the European defence programme FCAS, from discovery through to production rollout
- Delivered autonomous feature sets end-to-end: user stories, acceptance criteria, implementation support, testing, and GitOps-based rollout via ArgoCD
- Improved data quality through filtering, tagging, reproducible transformation pipelines, and audit trails with full data provenance (SafeAI, ExplainableAI concepts)
- Led a cross-functional team across architect, data scientists, backend and frontend engineers, and QA with multinational stakeholders
- Reduced deployment times from several days to 20–40 minutes by introducing GitOps with ArgoCD; independently unblocked infrastructure-level OpenShift issues that had stalled engineering for over a year

FEDERAL EMPLOYMENT AGENCY (BUNDESAGENTUR FÜR ARBEIT) 2023, 2024

Product Manager, Public Data Analytics

- Drove product development on the Arbeitsmarktmonitor, an analytics platform for regional labour market and demographic data
- Delivered labour shortage analytics, flexible regional comparison tooling, and modernised admin interfaces
- Introduced Dagster for automated ETL and indicator generation; standardised CI/CD with full test integration at build time
- Migrated legacy systems to Kubernetes and managed major platform upgrades
- Served as Deputy IT Security Officer with responsibility for compliance, governance, and data lineage

EMIL GROUP 2022, 2022

Product Manager, InsurTech SaaS Platform

- Delivered a B2B InsurTech platform reducing insurance product time-to-market from over one year to under one week, serving primary insurers, reinsurers, MGAs, brokers, and underwriters
- Designed and delivered a Claims Center module from concept through production rollout
- Re-established sprint discipline by splitting a bloated 17-person team into focused Scrum units, reintroducing velocity tracking and structured refinement
- Regained client trust in a prototype-stage environment through structured communication and weekly product discovery workshops
- Reduced senior developer attrition from near-total to a single departure through improved team clarity and ownership structures

CLINET PLATFORMS 2021, 2022

Product Owner, Mobile Healthcare App

- Delivered a mobile-first pilot app for hospitals across iOS and Android covering digital anamnesis, meal and therapy plans, transport coordination, chat, and patient document storage
- Integrated CGM interface into hospital information system workflows under tight deadlines
- Built stakeholder trust through transparent communication of technical complexity across legacy APIs and external databases

§ FOUNDER & EARLY-STAGE EXPERIENCE

- **Qcrypt AG (2016–2018), Product Owner:** Quantum-secure encryption product built on one-time-pad keys and true random number generator hardware; USP of setup under two minutes with no technical knowledge required. Also served as deputy IT security officer.
- **Kvitt Payment Solutions (2013–2018), Founder & CFO:** Group-based P2P payment product on a P2P backbone, including PSP integration, full UX/UI design, and mobile app prototype. Later acquired and operated by Sparkasse.
- **Smart Soil Technologies (2016–2021), Co-Founder/CDO/CFO/Project Lead:** Indoor vertical farming technology startup; established digital tooling, financial planning, investor relations, and HR from the ground up.
- **Bundesdruckerei (2021), Product Owner:** COVID-19 data platform for RKI projects DIM and DESH; secure data structures and identity management.

- **AOK Health Insurance ITSCare (2019–2020), Proxy Product Owner:** Migrated internal webshop to Shopware with complex catalogue and API integration and customisable approval policies.

§ LANGUAGES

German, native. English, C2.

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)