

Philipp Hilbert. *AI-native Product Manager for enterprise workflow transformation.*

Positioning, ten years translating complex operational workflows into production software, now building AI-native products where the requirement is to reinvent the work itself.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear CHAPTR team,

I have spent the past decade doing one thing in different settings: taking a messy, partially understood operational workflow and building the software product that makes it work at scale. At Rohde & Schwarz I did it inside a European defence programme, building a central AI data management platform for multimodal datasets from scratch. At Emil Group I did it inside B2B insurance, delivering a Claims Center module while simultaneously re-establishing sprint discipline for a 17-person team that had lost its delivery rhythm. At the Federal Employment Agency I did it for public labour market analytics, introducing Dagster-based ETL pipelines and modernising the platform to Kubernetes. The context changes; the core work does not.

CHAPTR's mandate is precisely that work, applied to the organisations that need it most: large enterprises trying to use generative AI to reshape daily workflows rather than bolt AI onto legacy processes. The job description asks for someone who can lead Product Discovery through Design Thinking, translate user needs into validated solutions, and drive iterations based on real feedback. That is my operating model. I run structured discovery workshops, write testable hypotheses, prototype fast in Lovable or Framer, and bring user interview evidence back to the team before engineering commits a single sprint. I have done this in regulated environments with demanding external stakeholders, which means I also know how to hold a product narrative steady when pressure comes from multiple directions at once.

I am genuinely AI-native, not as a talking point but as a daily working practice. I use Claude and GPT throughout the requirements and spec process, Cursor for in-repo review, and I have built a personal multi-agent orchestration framework in my own time to understand where these tools break under load. That background means I can have a credible technical conversation with the AI engineers on your team, push back on scope when the model is being asked to do something it cannot reliably do, and also protect the user experience when AI capability outpaces usability.

The specific pull of CHAPTR for me is the combination of startup execution pace inside a group that has real enterprise distribution. Springer Nature and Macmillan give the products a path to large-organisation customers; the startup structure means a PM can still make decisions at speed. Reporting directly to the Director of Product removes abstraction layers and puts accountability exactly where it belongs. I am a direct, low-ego operator who produces more in flat structures than in hierarchical ones, and the role description makes clear that is the culture you are building.

I would be glad to walk through any of the product decisions I have made in the past three years in a call. I can also show working prototypes and the agent framework I have built independently, which gives a clearer picture of how I actually work than a CV can.

Best regards, *Philipp Hilbert*
