Philipp Hilbert. *AI Product Manager driving investment-grade AI from data layer to decision.*

Positioning, Product manager with 13 years in regulated, data-intensive environments, now building production AI systems end-to-end across data pipelines, model layers, and user-facing workflows.

BERLIN, GERMANY · HILBERT@TRUE-NORTH.BERLIN · TRUE-NORTH.BERLIN

LinkedIn · GitHub · www.true-north.berlin · YouTube (EN sample)

Dear Hiring Team at LIQUIDITY,

Private credit is one of the few domains where a wrong AI output does not produce a bad UX: it produces a bad investment decision. That consequence-bearing quality is what drew me to your posting. I have spent the last two years building and owning production AI systems, for the FCAS European defence programme and for Germany's Federal Employment Agency, and in both cases the hardest problem was not prompting or model selection. It was ensuring that the outputs were accurate, auditable, and trusted by the analysts who depended on them. That is the same problem your AI Product Manager role is asking someone to solve.

At Rohde & Schwarz I led the internal AI and data platform for FCAS, managing 15+ engineers and data scientists through a full GitOps delivery cycle. The platform covered data ingestion, model layers, and developer-facing tooling. I introduced ArgoCD-based rollouts that cut deployment time from days to 30 minutes and made the platform hotfix-capable, which matters when the users are defence programme leads who cannot wait for a scheduled release window. I also instrumented competing product approaches empirically: when a UI-first workflow recorded exactly zero clicks against a developer-code-first alternative, the data settled the strategy debate in one sprint.

At the Federal Employment Agency I owned a national labour-market analytics platform for policymakers and researchers, served as Deputy IT Security Officer, and rebuilt the ETL layer with Dagster and dbt, cutting the time per quarterly data update by 75%. The platform ingested heterogeneous structured and unstructured data sources under GDPR constraints, which forced me to think carefully about data provenance, transformation auditability, and explainability: concepts that translate directly to the financial data mapping and risk scoring systems your posting describes.

Beyond employment, I have built a two-sided B2B marketplace from zero using LLM agents for funnel qualification and matching, and I actively trade across Hyperliquid, Bitget, and DEX venues, working with liquidation heatmaps and market microstructure analytics. This gives me a practitioner's understanding of financial data, not as rows in a schema, but as signals with timing, latency, and adversarial noise properties that matter when you are building AI systems that inform capital allocation decisions.

The intersection of multi-agent workflow design, structured financial data, and production-grade AI quality is where I have been concentrating my work. Liquidity is building something at that intersection at a scale and pace I find compelling. I would welcome the opportunity to discuss how my background fits the specific workflows you are prioritising.

Best regards, *Philipp Hilbert*
