

Philipp Hilbert. *AI-native Product Manager for enterprise SaaS.*

Positioning, ten years shipping B2B SaaS products end-to-end, from regulated insurance platforms to defence-grade AI data pipelines, now bringing that breadth to HR tech.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear hiring team at métier,

Strategic HR is one of the few SaaS categories where data sensitivity, regulatory exposure, and genuine business-criticality collide inside a single product. That combination is exactly where the product discipline I have built over the past decade applies most directly. I have shipped B2B SaaS platforms for insurers, reinsurers and brokers at EMIL Group, run a national analytics platform at Bundesagentur für Arbeit under recurring audit cycles, and owned the AI data backbone module for a European defence programme at Rohde & Schwarz. Across all three, the constant was translating ambiguous, high-stakes requirements into working software under tight compliance constraints.

The posting specifically highlights RAG pipelines, vector databases, and LLM orchestration as differentiating experience. At Rohde & Schwarz I owned end-to-end delivery of reproducible ML training pipelines with full data provenance and audit trails, reducing deployment times from several days to under forty minutes and unblocking an infrastructure problem the engineering team had not resolved in over a year. I worked directly with data scientists and architects to make trade-off decisions on pipeline design, not as a relay for stakeholder requests but as a participant in the technical conversation. That is the working mode your posting describes, and it is the one I default to.

At EMIL Group I built a claims module from zero, with no specification and no internal precedent. I ran discovery directly with a pilot customer's claims specialists, instrumented the claim flow so that churn per process step became the leading prioritisation metric, and restructured a seventeen-person delivery organisation that had stalled. Velocity recovered, senior engineers stopped leaving, and client trust returned. The parallel to an HR SaaS product is direct: the domain expert user base, the compliance overlay, and the need to earn engineering alignment rather than assume it are structurally identical.

My AI-native working mode changes the economics of product work. Using Claude Code, GPT, Lovable, and a multi-agent framework I built and operate myself, I routinely deliver output that previously required a four-person team. For a GenAI SaaS PM role, this matters in two ways: I can prototype and validate faster than a team waiting on engineering cycles, and I understand from the inside what it actually costs to build, deploy, and monitor an LLM-backed feature, which makes the accuracy-latency-cost trade-offs in the posting something I navigate from experience rather than theory.

I am available at four weeks' notice and happy to discuss the scope, team structure, and roadmap in detail.

Best regards, *Philipp Hilbert*
