

Philipp Hilbert. *Decision engine builder who ships.*

Positioning, ten years translating complex operational workflows into scalable software systems, now ready to own Taktile's core decisioning layer. ≤ 25 words.

BERLIN, GERMANY · [HILBERT@TRUE-NORTH.BERLIN](mailto:hilbert@true-north.berlin) · [TRUE-NORTH.BERLIN](https://true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

§ WHY TAKTILE

Taktile sits at exactly the intersection I have spent a decade working toward: high-stakes, high-availability decisioning systems where data quality, auditability, and workflow clarity are not nice-to-haves but hard requirements. The job description names balance of innovation with a high reliability bar as a core expectation — that framing matches precisely how I approached the FCAS AI data platform at Rohde & Schwarz and the Arbeitsmarktmonitor at the Federal Employment Agency, both environments where a wrong decision in production had real downstream consequences. Working alongside customers like Mercury, Monzo, and Faire, and shaping what agent-first workflow authoring looks like in financial services, is the kind of scope I am looking to grow into next. *That is exactly the environment where my strengths compound.*

§ AI-NATIVE STACK

Daily tools, not buzzwords.

- **Claude**, spec drafting, edge-case generation, code review
- **GPT**, research synthesis and structured requirement generation
- **GitHub Copilot**, in-repo implementation support
- **Cursor / VSCode**, AI-assisted development and refactoring
- **Lovable**, rapid front-end prototyping
- **Framer**, design-heavy mockups and interactive flows
- **n8n**, multi-agent orchestration and workflow automation
- **Gamma.app**, stakeholder-facing product narratives

§ TECHNICAL ENVIRONMENT

Shipped with, not just listed.

Production work delivered across Kotlin (Spring Boot), Python, Django, JavaScript, PostgreSQL, MongoDB, dbt, Dagster, DuckDB, Redpanda, Kubernetes, ArgoCD, Argo Workflows, Helm, OpenShift, Jenkins, REST APIs, AngularJS, Ionic, MicroStrategy, and GitOps toolchains. SQL proficiency used daily for ETL pipeline design, indicator generation, and data lineage work at the Federal Employment Agency and Rohde & Schwarz.

§ HOW I WOULD WORK AT TAKTILE

First weeks, concrete.

WEEK 1

Listen, map, find the gaps.

Shadow customer-facing teams on active decisioning use cases – lending, AML, fraud – and map the current decision engine's workflow authoring surface against the highest-friction points customers and engineers name most often.

WEEK 2

Slice and validate.

Take the two or three friction points that appear across multiple customer segments, write sharp problem statements with acceptance criteria, and pressure-test them with engineering on feasibility before committing any roadmap space.

WEEK 3+

Ship, measure, iterate.

Drive the first concrete improvement through the full delivery loop – user story, implementation support, rollout, metric baseline – and establish the review cadence that keeps reliability and innovation in balance going forward.

I COVER LETTER

Taktile's Decision Engine sits at the operational core of every lending approval, fraud check, and AML workflow your customers run. When that layer has gaps – in workflow clarity, in auditability, in reliability under load – the downstream cost is not a product complaint but a

rejected borrower, a missed fraud signal, or a compliance finding. I have spent nearly a decade working on exactly that category of problem, and I want to bring that experience to bear on Taktile's decisioning platform.

At Rohde & Schwarz, I owned the AI data management platform for the FCAS European defence programme. The core challenge there was not building features – it was building a system that other systems could trust: reproducible data transformation pipelines, full data provenance and audit trails, and SafeAI evaluation concepts that made model decisions legible to multinational stakeholders. I delivered this end-to-end, from user stories through implementation support and GitOps-based rollout, and I unblocked infrastructure-level blockers in OpenShift independently – one week to deliver a feature that had been stuck in the backlog for over a year. That combination of product ownership and technical depth is what I would bring to the Decision Engine.

At the Federal Employment Agency, I took ownership of the Arbeitsmarktmonitor – an analytics platform for regional labour market data used by policymakers and researchers. I introduced Dagster for automated ETL, migrated legacy systems to Kubernetes, standardised CI/CD with full test integration, and served as Deputy IT Security Officer with accountability for compliance and data lineage. This role sharpened my understanding of what it means to own a platform that underpins critical decisions at scale, where data quality and system reliability are non-negotiable.

Earlier, at Emil Group, I stepped into a prototype-stage InsurTech B2B SaaS environment where delivery had broken down, client trust was eroding, and a 17-person team was moving without ownership or structure. I split the team into smaller Scrum units, reintroduced velocity tracking and structured refinement, and led weekly product workshops that rebuilt client confidence. Reducing senior developer attrition from near-total to a single departure was a direct consequence of restoring clarity and discipline. That experience taught me that platform reliability is as much a product and process problem as a technical one – and that the PM is accountable for both.

I am based in Berlin, fluent in German and English, and ready to operate with the autonomy and ownership your role description outlines. The opportunity to define what agent-first decisioning looks like in financial services – working alongside the engineering team, close to customers like Mercury and Monzo, and grounded in real production constraints – is exactly the scope I am ready for.

Best regards, *Philipp Hilbert*

§ BEARING

Product Manager with nearly ten years of experience owning complex, high-availability software systems across regulated and data-intensive environments. Background spans AI and data platforms, InsurTech B2B SaaS, public analytics infrastructure, and founder-level company building. Comfortable translating ambiguous operational requirements into structured delivery, leading cross-functional teams without formal authority, and maintaining reliability bars while driving roadmap innovation. Strongest in environments where data quality, auditability, and workflow clarity carry real downstream consequences.

§ CORE STRENGTHS

- Translating complex operational workflows into dev-ready requirements
 - Platform ownership in high-availability, regulated environments
 - Data pipeline product design with auditability and provenance
 - Cross-functional alignment across engineering, data science, and compliance
 - Backlog management, sprint discipline, and delivery velocity restoration
 - SQL proficiency and data-driven prioritisation
 - AI-native product execution and rapid prototyping
 - Technical fluency across modern cloud-native stacks
 - Founder-level pragmatism: low ego, high agency, high ownership
-

§ PROFESSIONAL EXPERIENCE

ROHDE & SCHWARZ 2024, 2025

Product Manager, AI Data Transformation

- Designed and delivered reproducible data transformation pipelines with full audit trails and data provenance for a multinational European defence AI platform (FCAS)
- Enabled SafeAI and ExplainableAI evaluation concepts to make model decisions legible and auditable for non-technical stakeholders
- Reduced deployment times from several days to 20–40 minutes by introducing GitOps with ArgoCD; unblocked OpenShift infrastructure blockers independently – delivering in one week what had been outstanding for over a year
- Delivered autonomous feature sets end-to-end: user stories, acceptance criteria, implementation support, testing, and rollout

- Led a cross-functional team including architect, data scientists, backend and frontend engineers, and QA

Tech: Kotlin (Spring Boot), Python, PostgreSQL, MongoDB, Redpanda, Kubernetes, Argo Workflows, ArgoCD, OpenShift

FEDERAL EMPLOYMENT AGENCY (BUNDESAGENTUR FÜR ARBEIT) 2023, 2024

Product Manager, Public Data Analytics

- Owned the Arbeitsmarktmonitor analytics platform for regional labour market data used at policy and research level
- Introduced Dagster for automated ETL and indicator generation; maintained dbt and DuckDB for data modelling and lineage
- Delivered labour shortage analytics, flexible regional comparison features, and modernised admin tooling
- Migrated legacy systems to Kubernetes and standardised CI/CD with full test integration
- Served as Deputy IT Security Officer with accountability for compliance, governance, and data lineage

Tech: Python, Django, JavaScript, PostgreSQL, dbt, Dagster, DuckDB, MicroStrategy, Kubernetes, Jenkins

EMIL GROUP 2022, 2022

Product Manager, InsurTech SaaS Platform

- Owned product delivery for a B2B InsurTech platform reducing insurance product time-to-market from over one year to under one week, serving primary insurers, reinsurers, MGAs, brokers, and underwriters
- Designed and delivered a Claims Center module from concept to production rollout
- Re-established sprint discipline and delivery stability: split a bloated 17-person team into focused Scrum units, introduced ownership and deputy concepts, reintroduced velocity tracking
- Rebuilt client trust in a prototype-stage environment through structured communication and weekly product workshops
- Reduced senior developer attrition from near-total to a single departure

Tech: Low-code platform, Java, HTML/CSS, Kubernetes, REST APIs

CLINET PLATFORMS 2021, 2022

Product Owner, Mobile Healthcare App

- Delivered core flows for digital anamnesis, meal and therapy plans, transport, chat, and patient document storage for a mobile-first hospital app (iOS and Android)
- Integrated CGM interface into hospital information system workflows, managing complex legacy API dependencies transparently with stakeholders
- Led a small team under tight deadlines with strict focus on highest-leverage features

Tech: AngularJS, Ionic, Python, Kubernetes, CGM API integration

BUNDESDRUCKEREI 2021, 2021

Product Owner, COVID-19 Data Platform

- Owned delivery on RKI projects DIM and DESH: secure data structures and identity management for national health data infrastructure
-

AOK HEALTH INSURANCE – ITSCARE 2019, 2020

Proxy Product Owner

- Led migration of internal webshop to Shopware, including complex product catalogue and API integration and configurable approval policy workflows
-

§ FOUNDER & EARLY-STAGE EXPERIENCE

QCRYPT AG 2016, 2018

Product Owner & Deputy IT Security Officer

- Owned product architecture for a quantum-secure encryption product (one-time-pad keys, true-random-number-generator hardware) targeting enterprise and government customers
- Designed for setup under two minutes with no required technical knowledge; served as deputy IT security officer

KVITT PAYMENT SOLUTIONS 2013, 2018

Founder & CFO

- Founded a group-based P2P payments product on a P2P backbone; owned full UX/UI design, PSP integration, and mobile app prototype
- Product later acquired and operated by Sparkasse

SMART SOIL TECHNOLOGIES 2016, 2021

Co-Founder, CDO & CFO

- Established digital tooling, financial planning, and investor relations for an AgTech venture in intelligent substrate regulation for vertical indoor farming

§ LANGUAGES

German, native. English, C2.

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

TRUE NORTH · APPLICATION PACKAGE · TAKTILE

