

Philipp Hilbert. *AI Product Manager for document intelligence and extraction.*

Positioning. I have shipped AI/ML products in regulated B2B environments where a wrong number costs real money, and I know how to turn model metrics into decisions engineers and executives can act on.

BERLIN, GERMANY · HILBERT@TRUE-NORTH.BERLIN · [TRUE-NORTH.BERLIN](https://www.true-north.berlin)

[LinkedIn](#) · [GitHub](#) · www.true-north.berlin · [YouTube \(EN sample\)](#)

Dear K1X team,

I have been a product manager on AI and data products in regulated B2B environments for several years, most recently at Rohde & Schwarz where I owned the internal AI and data platform for the FCAS European programme. The work involved exactly the trade-offs your job description names: precision versus coverage, evaluation set design, cost per inference, and how to explain model failures to stakeholders who care about outcomes, not model internals.

Your posting is specific about what makes this role hard. You need someone who can own the accuracy definition end-to-end, from per-field precision to straight-through rate to the user corrections that show up in Client Success tickets. At the Federal Employment Agency I built the ETL and data pipeline that fed a national labour-market analytics platform. Accuracy in that context was not abstract: a wrong indicator number in a published dataset was a policy error. I introduced Dagster and dbt to make quarterly data updates repeatable and auditable, and I integrated test automation into the Jenkins pipeline so QA could verify data quality without hand-holding. That is the kind of unglamorous, high-stakes data discipline your accuracy report work requires.

On prioritisation and roadmap planning, I have run RICE and opportunity scoring on backlogs at EMIL Group (18-person delivery organisation, B2B SaaS fintech) and at CLINET Platforms (healthcare, five-person team). At CLINET I cut the roadmap to the top 10% of features against agreed success measures, and user acceptance rose more than fivefold. I know how to size a roadmap against real team capacity and how to sequence dependencies when the September-to-November filing peak means you cannot afford a late-breaking scope change.

The tax domain is new to me, and I am comfortable saying so. I have worked in defence, public sector, insurance, healthcare, and payments, all domains where I had to learn the failure modes fast. What I bring is the framework for doing that quickly: I find the people who know where the footnotes are and I ask them the right questions.

I would welcome a conversation about how my background fits what you are building.

Best regards, *Philipp Hilbert*
